

Cannabis Sales Revenue Report

Nuutti Barron - COM SCI M148

UID: 805205422

Executive Summary

For this project we analyze a large dataset of licensed cannabis sales in the state of California in order to determine factors that potentially affect sales, and to develop models meant to predict future sales. The first step is to merge together data from multiple sources, as well as combine time series data with product and brand information. Following this step we gather basic statistics, such as mean, median, standard deviation, and visualized distributions of all the features we analyze. This provides a general sense of what the data we are working with looks like. We also calculate correlations between all features in hopes of identifying factors correlated with monthly sales, as these can be adjusted to improve sales going forward. The features that we discovered to be most strongly correlated with monthly sales revenue are outlined in the following table:

Positive Correlation	Negative Correlation
Offer products with Pax filters	Mean mg THC across all products
Offer product types other than inhalables and ingestibles	Average Retail Price
Offering CBD infused products	
The number of products offered	
Price range in products offered	
Number of products with a mood effect	
Number of products that are flavored	

Note: Factors that we found to be correlated with sales, but do not provide useful information (such as revenue in the previous month), are excluded from this table

Following this step we trained Linear Regression, K-Nearest Neighbor (KNN), Random Forest, and Neural Network models in order to predict future sales. We implement hyperparameter optimization to optimize each model, and cross validation in order to build confidence in the accuracy of each model. Attempted principal component analysis yielded decreased performance metrics across each model, so we do not use this technique to alter our dataset when training and testing each model. Below we provide the mean absolute error (MAE), the average magnitude of error in predicted sales, and mean absolute percentage error (MAPE), the average percent error, of each model we built:

Model	MAE
Linear Regression	83673.16
KNN Regression	80769.39
Random Forest Regression	78643.26
Neural Network Regression	83010.92

Introduction and Background

In this project we analyze cannabis sales data from a wide variety of companies in the state of California collected over the last few years. Our goal is to identify potential factors that influence sales, as well as develop computational, data-driven models to accurately predict future sales. There are various reasons that make meaningful data analysis and the ability to build accurate predictive models complicated. One reason is the legal status of cannabis, and the varying regulatory status of cannabis across the United States. The legal status varies from state to state, meaning the supply chain for the industry is nonuniform, which affects the type of products sold, the demand, and other factors in the cannabis industry. This means that meaningful data analysis in one region will likely not translate to another. Our data is focused on licensed sale of cannabis in California (where Cannabis is legal for medical and recreational use), which alleviates this complexity, but it nevertheless affects our dataset. Another factor is the presence of new forms of cannabis products (other than inhalable and ingestible). The emergence of CBD based products, Pax filter products, flavored, and mood specific products is

an example of this. Legal cannabis sales is an industry that is still emerging, making it generally more unpredictable.

In order to identify potential factors that influence sales, we analyze the dataset provided by Cookies. The provided dataset includes time series data for cannabis sales, meaning it provides timestamped information about total sales, total units sold, and what types of products were sold, among other information. It also contains more detailed brand information about individual brands and products. In order to analyze the dataset, the data from different sources is merged to a common data frame to be further analyzed. This step involves combining and augmenting time series data with undated product data. From the combined dataframe, we obtain various statistics, such as the mean, median, and standard deviation, for all of the variables that we consider, providing a sense of what the data looks like. For example, we extract that the average monthly sales across all brands is \$409,372.90 with a standard deviation of \$1,596,024 , meaning that there is a lot of variability in total sales, and that it does not resemble a normal distribution. We also visualize the distribution of the data to get a sense of what we are working with. More interestingly, we identify correlations between the variables, or features, that we consider. Most important are the features that have a strong correlation (negative or positive) with total sales, as these can be changed in hopes of boosting cannabis sales. Some factors, such as the sales from the previous month, are strongly correlated, but trivial, in that it does not provide any useful information. Variables found to have a non-negligible correlation with sales are outlined in the executive summary above.

Using the features we found to be correlated with sales, we train models intended to predict future sales. First we augment the dataset to make it suitable for model training. This includes scaling nonbinary features, encoding categorical features, and imputing the data to get rid of unknown, or null, values. We implement a pipeline for all of this, making the data augmentation easily reproducible. Following these steps, we build a simple Linear Regression, a K-Nearest Neighbor Regression, a Random Forest Regression, and a Neural Network Regression model in order to predict future monthly sales. In order to maximize the performance of each model, we implement hyperparameter optimization, and in order to validate the performance of each model, we extract cross fold validation metrics. We test the effect principal component analysis (to reduce the dimensionality of the dataset) has on our model performances, but it

resulted in decreased model performance, so for our finalized models we do not do this to the dataset.

Methodology

In order to develop a dataset we can analyze and train on future models, we must merge data from multiple data files. An added layer of complexity is added since some of the data files contain time series data and some do not. In order to make one unified dataset, we remove the dates of sales made and implement a feature for previous months sales, as well as a feature for the brand's sales over a three month rolling average. We can then include the variables from the undated data. For example, if a brand indicates that they sell CBD based products, we add a feature to indicate that the brand sells CBD products. We follow this methodology for the important features in the undated data (e.g. whether the brand sells flavored products, mood directed product, Pax filters, ingestible products, etc). Although this step may not be necessary for basic data exploration and analysis, this step is crucial to build predictive models using time series data.


The dataset is further altered for a few reasons. First of all, categorical features must be encoded, so we implement our own binary encoding for categorical features that is similar to one hot encoding (in our features, more than one categorical value is possible, one hot encoding alone would not suffice). Some features also have greater variance in their data points than others, which can affect model training, and bias the importance of these features. In order to counteract this, we apply a standard scalar to features that contain nonbinary, continuous data. Finally, we must impute unknown values in the dataset. Different features are imputed in different ways. For binary features that indicate if a brand sells a certain type of product, we assume that the brand does not sell this type of product if it is not indicated, so we set the unknown values in these columns to 0. For scalar features, we assign the median value of the known data. The reasoning behind this is that there are few unknown values relative to the total amount of data in these columns, but these features are not normally distributed, so applying the mean would likely skew models trained on this data more than imputing the median. These steps are all implemented in a pipeline fashion, so these steps are easily reproducible.

Following the implementation of a Linear Regression model, we implemented a random forest regression to test the efficacy of an ensemble regression method on our dataset. As opposed to linear regression, this model has hyperparameters that we tune in order to optimize the model. The hyperparameters that we optimized are the maximum tree depth, and the number of trees in the random forest. We actually found improved performance metrics as opposed to simple linear regression using this model.


In order to validate the results of our model, we apply K-fold cross validation. In doing this, we split the dataset into 10 train-test splits, train the random forest model on each split, retrieve performance metrics on the current split, and average these performance metrics. This provides a greater level of confidence in the validity of the models we build and train. The variance in performance metrics from our original models was essentially negligible, providing greater confidence in their validity and reproducibility.

Results


After merging all of our time series and brand detail data into one dataframe, we analyze the basic statistics for the features in the dataset. The following displays distributions for a subset of the features that we consider.




Price Range




Pax Filter Product Offered




Inhalable Products Offered




Ingestible Products Offered




Number of Flavored Products Offered



Number of Mood Specific Products Offered



We also obtain a heatmap of correlations between important variables in our unified dataframe. The features correlated (positively or negatively) with 'Total Sales (\$)' are the features we deem the most important, as this is the value the models we build will try to predict.



Note: Features 1-12 here indicate the month of the year of the current entry. I suspected that sales may generally increase in certain months and decrease in others, but this heat map indicates otherwise


After running basic statistical analysis on our data, we build a linear regression model with 17 features used to predict future sales. Our linear regression reported the following metrics on the testing set:

$$R^2 = 0.9776, MAE = 89337.48, MSE = 71251360719.15, RMSE = 266929.51$$

Coefficients for Each Feature

Feature 0 - Last_Month Score: -16726.93808
Feature 1 - 3_Month_Avg Score: -1252.40326
Feature 2 - Last_Year Score: -9613.25731
Feature 3 - ARP Score: 41837.51016
Feature 4 - Pax_Filer Score: 10005.31014
Feature 5 - Inhaleables Score: 2039.77175
Feature 6 - Ingestibles Score: 966.01786
Feature 7 - Other_Cannabis Score: -4060.54556
Feature 8 - Topicals Score: -6173.71553
Feature 9 - Max_THC Score: 13567.07529
Feature 10 - Sell_CBD Score: 12725.10978
Feature 11 - Product_Count Score: 27512.33976
Feature 12 - Price_Range Score: -10621.30227
Feature 13 - Flavored_Count Score: 18818.28732
Feature 14 - Mood_Count Score: 245711.08046
Feature 15 - Mean_THC Score: 1453753.73643
Feature 16 - MTHC_Ing_Cross Score: 81676.23301
Feature 17 - New_Product Score: -149.58918

Bar graph of the Feature Coefficient



Other statistical metrics for our linear regression

feature	Coef	std err	t	P> t	[0.025	0.975]
const	1.616e+05	1.07e+04	15.154	0.000	1.41e+05	1.82e+05
x1	-1.673e+04	4277.383	-3.911	0.000	-2.51e+04	-8342.948
x2	-1252.4033	1.18e+04	-0.106	0.916	-2.45e+04	2.2e+04
x3	-9613.2573	2737.378	-3.512	0.000	-1.5e+04	-4247.792
x4	4.184e+04	3331.598	12.558	0.000	3.53e+04	4.84e+04
x5	1.001e+04	3166.037	3.160	0.002	3799.642	1.62e+04
x6	2039.7717	2988.370	0.683	0.495	-3817.656	7897.200
x7	966.0179	2587.289	0.373	0.709	-4105.261	6037.296
x8	-4060.5456	1.11e+04	-0.366	0.714	-2.58e+04	1.77e+04
x9	-6173.7155	1.53e+04	-0.403	0.687	-3.62e+04	2.38e+04
x10	1.357e+04	5931.178	2.287	0.022	1941.525	2.52e+04
x11	1.273e+04	8093.040	1.572	0.116	-3137.851	2.86e+04
x12	2.751e+04	1.2e+04	2.293	0.022	3995.970	5.1e+04
x13	-1.062e+04	1.15e+04	-0.923	0.356	-3.32e+04	1.19e+04
x14	1.882e+04	7134.866	2.638	0.008	4833.418	3.28e+04
x15	2.457e+05	1.03e+04	23.923	0.000	2.26e+05	2.66e+05
x16	1.454e+06	8437.734	172.292	0.000	1.44e+06	1.47e+06
x17	8.168e+04	9155.607	8.921	0.000	6.37e+04	9.96e+04
x18	-149.5892	4117.140	-0.036	0.971	-8219.491	7920.312

Using a p-value threshold of 0.05, we determine that the following variables are important in this linear regression: Previous months revenue, previous years revenue, average retail price, sells Pax filter products, Max mg THC in any product offered, number of products offered, number of flavored products offered, number of mood specific products offered, average THC across all products, and the cross feature of max THC offered and ingestible products offered.

The ensemble model we trained was a Random Forest. The following are the 10-fold cross validated metrics reported by the trained random forest on our dataset using the same features as our linear regression.

Random Forest Regression Performance Metrics

R^2	0.9484
MAE	87530.34
MSE	137768222451.09
RMSE	331004.82

In order to optimize our random forest model, we hyperparameter optimization using GridSearch. We found that setting the max depth to 15, and the number of estimators (or decision trees) to 50. With these parameters, we got the following cross validated performance metrics.

Random Forest Regression Performance Metrics After Hyperparameter Optimization

R^2	0.98
MAE	77355.78
MSE	42009717731.5782
RMSE	204962.7228

We see that the performance metrics all improved following hyperparameter tuning (R-squared increased, and MAE, MSE, and RMSE all decreased).

To test a few additional models, we implemented K-Nearest Neighbor and Neural Network regression models. We employ hyperparameter training as before, and display the optimal hyperparameters as well as predictive cross validated performance metrics for both models.

K Nearest Neighbor Regression
(neighbors = 3)

R^2	0.9493
MAE	80769.39
MSE	126561893173.77
RMSE	312018.23

Neural Network Regression
(number of hidden layers = 5, nodes in each layer = 5)

R^2	0.9521
MAE	83010.92
MSE	120426132575.03
RMSE	289853.44

Based on our data analysis, the following features likely are important in determining cannabis sales: average retail price, selling Pax filter products, maximum mg THC in any product offered, number of products offered, number of flavored products offered, number of mood specific products offered, average THC across all products.

Discussion

The subset of variables that are correlated with cannabis sales we found are listed above. These results come from an analysis of licensed cannabis sales in the state of California, so the results may not extend out of this region. Due to the average error we find in our models, it means that they may not have a lot of predictive power (the mean absolute error is around \$80,000 for every model), but they can provide a good ballpark estimate for sales for a certain input.

Based on the analysis done in this project a number of things can be done to potentially increase sales. Offering CBD based products, Pax filtered products, and increasing the number of flavored and mood specific products are all indicators of increased sales. Increasing the price range of product units, and offering a wider array of products are also indicators of increased cannabis sales according to our analysis. Lowering the average retail price, and lowering the average amount of THC per product also could potentially increase monthly sales based on our data analysis.

For further analysis, cannabis sale data from outside of California should be collected and analyzed to increase the predictive power of models developed, and to allow models developed to work for states outside of California. A more detailed analysis on more detailed product descriptions could also be incorporated into the predictive models we produce.

Conclusion

In this project, we analyzed a dataset of licensed cannabis sales in the state of California over the period of a few years. We merged time series data with more detailed product information to form a dataset that we can train computational models on. We analyze the data by looking at the distribution of variables, as well as other basic statistical measures. We also calculate and extract correlations between all features in order to determine which features are useful in deciding monthly cannabis sales for a given brand or product. Following this, we further augment the dataset and prepare it for model training by scaling and imputing the data. Following this step, we train a Linear Regression, a K-Nearest Neighbor, a Random Forest, and a Neural Network regression model to predict future cannabis sales. These were found to have a

strong correlation with sales (as depicted by an R-squared value close to 1), but a high MAE score. Each model is validated using K-Fold cross validation, and optimized with hyperparameter optimization. We implemented principal component analysis on our dataset in order to reduce the data's dimensionality in an attempt to prevent overfitting, but found that the model performance decreased as a result of this.

proj3

December 4, 2021

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import os
import itertools
import random

import statistics
import seaborn as sns

%matplotlib inline
random.seed(148)
```

```
[2]: brandAvgRetPrice = pd.read_csv('BrandAverageRetailPrice.csv')
brandDetails = pd.read_csv('BrandDetails.csv')
brandTotalSales = pd.read_csv('BrandTotalSales.csv')
brandTotalUnits = pd.read_csv('BrandTotalUnits.csv')
```

```
[3]: brandAvgRetPrice.head(10)
```

```
[3]:
```

	Brands	Months	ARP	vs. Prior Period
0	#BlackSeries	08/2020	15.684913	NaN
1	#BlackSeries	09/2020	NaN	-1.000000
2	#BlackSeries	01/2021	13.611428	NaN
3	#BlackSeries	02/2021	11.873182	-0.127705
4	#BlackSeries	03/2021	NaN	-1.000000
5	101 Cannabis Co.	11/2019	34.066667	NaN
6	101 Cannabis Co.	12/2019	NaN	-1.000000
7	101 Cannabis Co.	01/2020	34.134929	NaN
8	101 Cannabis Co.	02/2020	29.091388	-0.147753
9	101 Cannabis Co.	03/2020	32.293498	0.110071

```
[4]: brandAvgRetPrice.describe()
```

```
[4]:
```

	ARP	vs. Prior Period
count	25279.000000	24499.000000
mean	22.679732	-0.065028

std	19.802724	0.388923
min	0.000000	-1.000000
25%	10.512827	-0.088073
50%	17.033051	-0.011649
75%	31.505612	0.045232
max	700.874984	12.645741

```
[5]: brandAvgRetPrice.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27211 entries, 0 to 27210
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brands                 27211 non-null  object
1   Months                 27211 non-null  object
2   ARP                    25279 non-null  float64
3   vs. Prior Period      24499 non-null  float64
dtypes: float64(2), object(2)
memory usage: 850.5+ KB
```

```
[6]: brandDetails.head(50)
```

```
[6]:
```

	State	Channel	Category L1	Category L2	Category L3 \
0	California	Licensed	Inhaleables	Flower	Hybrid
1	California	Licensed	Inhaleables	Flower	Hybrid
2	California	Licensed	Inhaleables	Flower	Sativa Dominant
3	California	Licensed	Inhaleables	Flower	Sativa Dominant
4	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
5	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
6	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
7	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
8	California	Licensed	Inhaleables	Pre-Rolled	Infused Pre-Rolled
9	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
10	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
11	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
12	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
13	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
14	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
15	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
16	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
17	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
18	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
19	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
20	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
21	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
22	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates
23	California	Licensed	Inhaleables	Concentrates	Dabbable Concentrates

24	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
25	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
26	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
27	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
28	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
29	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
30	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
31	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
32	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
33	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
34	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
35	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
36	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
37	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
38	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
39	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
40	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
41	California	Licensed	Inhaleables	Pre-Rolled	Infused	Pre-Rolled
42	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
43	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
44	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
45	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
46	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
47	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
48	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates
49	California	Licensed	Inhaleables	Concentrates	Dabbable	Concentrates

	Category L4	Category L5	Brand \
0	NaN	NaN	#BlackSeries
1	NaN	NaN	#BlackSeries
2	NaN	NaN	#BlackSeries
3	NaN	NaN	#BlackSeries
4	Wax	NaN	101 Cannabis Co.
5	Wax	NaN	101 Cannabis Co.
6	Wax	NaN	101 Cannabis Co.
7	Wax	NaN	101 Cannabis Co.
8	NaN	NaN	101 Cannabis Co.
9	Wax	NaN	101 Cannabis Co.
10	Wax	NaN	101 Cannabis Co.
11	Wax	NaN	101 Cannabis Co.
12	Wax	NaN	101 Cannabis Co.
13	Wax	NaN	101 Cannabis Co.
14	Wax	NaN	101 Cannabis Co.
15	Wax	NaN	101 Cannabis Co.
16	Wax	NaN	101 Cannabis Co.
17	Wax	NaN	101 Cannabis Co.
18	Wax	NaN	101 Cannabis Co.

19	Wax	NaN	101 Cannabis Co.
20	Wax	NaN	101 Cannabis Co.
21	Wax	NaN	101 Cannabis Co.
22	Wax	NaN	101 Cannabis Co.
23	Wax	NaN	101 Cannabis Co.
24	Wax	NaN	101 Cannabis Co.
25	Wax	NaN	101 Cannabis Co.
26	Wax	NaN	101 Cannabis Co.
27	Wax	NaN	101 Cannabis Co.
28	Wax	NaN	101 Cannabis Co.
29	Wax	NaN	101 Cannabis Co.
30	Wax	NaN	101 Cannabis Co.
31	Wax	NaN	101 Cannabis Co.
32	Wax	NaN	101 Cannabis Co.
33	Wax	NaN	101 Cannabis Co.
34	Wax	NaN	101 Cannabis Co.
35	Wax	NaN	101 Cannabis Co.
36	Wax	NaN	101 Cannabis Co.
37	Wax	NaN	101 Cannabis Co.
38	Wax	NaN	101 Cannabis Co.
39	Wax	NaN	101 Cannabis Co.
40	Live Resin	NaN	101 Cannabis Co.
41	NaN	NaN	101 Cannabis Co.
42	Wax	NaN	101 Cannabis Co.
43	Wax	NaN	101 Cannabis Co.
44	Wax	NaN	101 Cannabis Co.
45	Wax	NaN	101 Cannabis Co.
46	Wax	NaN	101 Cannabis Co.
47	Wax	NaN	101 Cannabis Co.
48	Wax	NaN	101 Cannabis Co.
49	Wax	NaN	101 Cannabis Co.

	Product Description	Total Sales (\$) \
0	#BlackSeries - Vanilla Frosting - Flower (Gram)	1,103.964857
1	#BlackSeries - Vanilla Frosting - Flower (Gram)	674.645211
2	#BlackSeries - Blueberry Slushy - Flower (Gram)	2,473.699102
3	#BlackSeries - Blueberry Slushy - Flower (Gram)	14,589.916417
4	101 Cannabis Co. - Afghan Kush - Wax	145.39627
5	101 Cannabis Co. - Skywalker OG - Wax	3,261.12486
6	101 Cannabis Co. - Skywalker OG - Wax	2,062.231412
7	101 Cannabis Co. - Indica Strain Blends - Wax	62.556665
8	101 Cannabis Co. - Hybrid Strain Blends - Infu...	1,309.279796
9	101 Cannabis Co. - Kosher Kush - Wax	556.738062
10	101 Cannabis Co. - Kosher Kush - Wax	1,316.637371
11	101 Cannabis Co. - Kosher Kush - Wax	9,225.549476000000
12	101 Cannabis Co. - Kosher Kush - Wax	3,019.525038000000
13	101 Cannabis Co. - Blood Orange - Wax	566.293122

14	101 Cannabis Co. - 3 Kings - Wax	6,261.743324
15	101 Cannabis Co. - 3 Kings - Wax	2,883.5042220000000
16	101 Cannabis Co. - 3 Kings - Wax	815.067402
17	101 Cannabis Co. - 3 Kings - Wax	407.66515
18	101 Cannabis Co. - 3 Kings - Wax	582.488434
19	101 Cannabis Co. - 3 Kings - Wax	826.2600400000000
20	101 Cannabis Co. - 3 Kings - Wax	105.603128
21	101 Cannabis Co. - Bubba Kush - Wax	145.39627
22	101 Cannabis Co. - Blood Orange - Wax	358.45266
23	101 Cannabis Co. - Sour Diesel - Wax	434.970731
24	101 Cannabis Co. - Blood Orange - Wax	86.466443
25	101 Cannabis Co. - MK Ultra - Wax	377.68802900000000
26	101 Cannabis Co. - Platinum OG - Wax	105.898451
27	101 Cannabis Co. - Platinum OG - Wax	101.972163
28	101 Cannabis Co. - Super Silver Haze - Wax	202.22070300000000
29	101 Cannabis Co. - Super Silver Haze - Wax	435.843645
30	101 Cannabis Co. - Sour Diesel - Wax	371.63322500000000
31	101 Cannabis Co. - Skywalker OG - Wax	1,527.8381580000000
32	101 Cannabis Co. - Sour Diesel - Wax	108.5442
33	101 Cannabis Co. - Skywalker OG - Wax	389.06757000000000
34	101 Cannabis Co. - Super Silver Haze - Wax	2,414.668334
35	101 Cannabis Co. - Durban Poison - Wax	920.0528850000000
36	101 Cannabis Co. - Lemonade - Wax	279.18337500000000
37	101 Cannabis Co. - Lemonade - Wax	7,766.681791000000
38	101 Cannabis Co. - Sundae Driver - Wax	292.619332
39	101 Cannabis Co. - Super Silver Haze - Wax	5,346.86861
40	101 Cannabis Co. - Zookies - Live Resin	1,196.454442
41	101 Cannabis Co. - Hybrid Strain Blends - Infu...	385.053811
42	101 Cannabis Co. - Durban Poison - Wax	2,121.358598
43	101 Cannabis Co. - Zookies - Wax	329.280835
44	101 Cannabis Co. - Lemon Skunk - Wax	472.14675800000000
45	101 Cannabis Co. - Durban Poison - Wax	182.056694
46	101 Cannabis Co. - Durban Poison - Wax	65.110869
47	101 Cannabis Co. - Lemon Skunk - Wax	1,325.24142
48	101 Cannabis Co. - Lemon Skunk - Wax	8,684.973093
49	101 Cannabis Co. - Lemon Skunk - Wax	1,683.402394

	... Total THC	Total CBD	Contains CBD	Pax Filter	Strain \	
0	...	0	0	THC Only	NaN	Vanilla Frosting
1	...	0	0	THC Only	NaN	Vanilla Frosting
2	...	0	0	THC Only	NaN	Blueberry Slushy
3	...	0	0	THC Only	NaN	Blueberry Slushy
4	...	0	0	THC Only	NaN	Afghan Kush
5	...	0	0	THC Only	NaN	Skywalker OG
6	...	0	0	THC Only	NaN	Skywalker OG
7	...	0	0	THC Only	NaN	Indica Strain Blends
8	...	0	0	THC Only	NaN	NaN

9	...	0	0	THC Only	NaN	Kosher Kush
10	...	0	0	THC Only	NaN	Kosher Kush
11	...	0	0	THC Only	NaN	Kosher Kush
12	...	0	0	THC Only	NaN	Kosher Kush
13	...	0	0	THC Only	NaN	Blood Orange
14	...	0	0	THC Only	NaN	3 Kings
15	...	0	0	THC Only	NaN	3 Kings
16	...	0	0	THC Only	NaN	3 Kings
17	...	0	0	THC Only	NaN	3 Kings
18	...	0	0	THC Only	NaN	3 Kings
19	...	0	0	THC Only	NaN	3 Kings
20	...	0	0	THC Only	NaN	3 Kings
21	...	0	0	THC Only	NaN	Bubba Kush
22	...	0	0	THC Only	NaN	Blood Orange
23	...	0	0	THC Only	NaN	Sour Diesel
24	...	0	0	THC Only	NaN	Blood Orange
25	...	0	0	THC Only	NaN	MK Ultra
26	...	0	0	THC Only	NaN	Platinum OG
27	...	0	0	THC Only	NaN	Platinum OG
28	...	0	0	THC Only	NaN	Super Silver Haze
29	...	0	0	THC Only	NaN	Super Silver Haze
30	...	0	0	THC Only	NaN	Sour Diesel
31	...	0	0	THC Only	NaN	Skywalker OG
32	...	0	0	THC Only	NaN	Sour Diesel
33	...	0	0	THC Only	NaN	Skywalker OG
34	...	0	0	THC Only	NaN	Super Silver Haze
35	...	0	0	THC Only	NaN	Durban Poison
36	...	0	0	THC Only	NaN	Lemonade
37	...	0	0	THC Only	NaN	Lemonade
38	...	0	0	THC Only	NaN	Sundae Driver
39	...	0	0	THC Only	NaN	Super Silver Haze
40	...	0	0	THC Only	NaN	Zookies
41	...	0	0	THC Only	NaN	NaN
42	...	0	0	THC Only	NaN	Durban Poison
43	...	0	0	THC Only	NaN	Zookies
44	...	0	0	THC Only	NaN	Lemon Skunk
45	...	0	0	THC Only	NaN	Durban Poison
46	...	0	0	THC Only	NaN	Durban Poison
47	...	0	0	THC Only	NaN	Lemon Skunk
48	...	0	0	THC Only	NaN	Lemon Skunk
49	...	0	0	THC Only	NaN	Lemon Skunk

	Is Flavored	Mood Effect	Generic Vendor	Generic Items \
0	NaN	Not Mood Specific	Non-Generic Vendors	Non-Generic Items
1	NaN	Not Mood Specific	Non-Generic Vendors	Non-Generic Items
2	NaN	Not Mood Specific	Non-Generic Vendors	Non-Generic Items
3	NaN	Not Mood Specific	Non-Generic Vendors	Non-Generic Items

	\$5 Price Increment
0	\$10.00 to \$14.99
1	\$15.00 to \$19.99
2	\$15.00 to \$19.99
3	\$10.00 to \$14.99
4	\$35.00 to \$39.99
5	\$30.00 to \$34.99
6	\$20.00 to \$24.99
7	\$10.00 to \$14.99
8	\$25.00 to \$29.99
9	\$30.00 to \$34.99
10	\$45.00 to \$49.99
11	\$35.00 to \$39.99
12	\$40.00 to \$44.99
13	\$25.00 to \$29.99
14	\$35.00 to \$39.99
15	\$40.00 to \$44.99
16	\$30.00 to \$34.99
17	\$50.00 to \$54.99
18	\$45.00 to \$49.99
19	\$25.00 to \$29.99
20	\$20.00 to \$24.99
21	\$35.00 to \$39.99
22	\$30.00 to \$34.99
23	\$35.00 to \$39.99
24	\$20.00 to \$24.99
25	\$45.00 to \$49.99
26	\$45.00 to \$49.99
27	\$40.00 to \$44.99
28	\$20.00 to \$24.99
29	\$35.00 to \$39.99
30	\$30.00 to \$34.99
31	\$25.00 to \$29.99
32	\$25.00 to \$29.99
33	\$15.00 to \$19.99
34	\$25.00 to \$29.99
35	\$25.00 to \$29.99
36	\$30.00 to \$34.99
37	\$35.00 to \$39.99
38	\$30.00 to \$34.99
39	\$30.00 to \$34.99
40	\$25.00 to \$29.99
41	\$20.00 to \$24.99
42	\$35.00 to \$39.99
43	\$35.00 to \$39.99
44	\$15.00 to \$19.99
45	\$20.00 to \$24.99

```

46  $15.00 to $19.99
47  $20.00 to $24.99
48  $35.00 to $39.99
49  $30.00 to $34.99

```

```
[50 rows x 25 columns]
```

```
[7]: brandDetails.describe()
```

```

[7]:
      count  ARP  Items Per Pack
count  144977.000000  144977.000000
mean      30.828439      1.938259
std       19.367580     17.294108
min        0.000000      0.000000
25%       16.407796      0.000000
50%       28.073823      0.000000
75%       41.781699      0.000000
max       874.800010     1000.000000

```

```
[8]: brandDetails.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 144977 entries, 0 to 144976
Data columns (total 25 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   State                 144977 non-null object
 1   Channel               144977 non-null object
 2   Category L1           144977 non-null object
 3   Category L2           144977 non-null object
 4   Category L3           144245 non-null object
 5   Category L4           102618 non-null object
 6   Category L5           50135 non-null  object
 7   Brand                 144977 non-null object
 8   Product Description    144977 non-null object
 9   Total Sales ($)       144977 non-null object
10   Total Units           144977 non-null object
11   ARP                   144977 non-null float64
12   Flavor                7807 non-null  object
13   Items Per Pack        144977 non-null int64
14   Item Weight           64454 non-null object
15   Total THC             144977 non-null object
16   Total CBD             144977 non-null object
17   Contains CBD          144977 non-null object
18   Pax Filter            44301 non-null object
19   Strain                115639 non-null object
20   Is Flavored           11287 non-null  object
21   Mood Effect           144977 non-null object

```

```

22 Generic Vendor      144977 non-null object
23 Generic Items       144977 non-null object
24 $5 Price Increment  144977 non-null object
dtypes: float64(1), int64(1), object(23)
memory usage: 27.7+ MB

```

```

[9]: print(len(brandDetails['Brand'].unique()))
pax_filter_options = brandDetails['Pax Filter'].unique()
print(str(pax_filter_options))
flavor_options = brandDetails['Is Flavored'].unique()
print(str(flavor_options))
mood_options = brandDetails['Mood Effect'].unique()
print(str(mood_options))

```

```

1123
[nan 'Not Pax' 'Pax']
[nan 'Not Flavored' 'Flavored']
['Not Mood Specific' 'Mood Specific']

```

```

[10]: brandTotalSales.head(10)

```

```

[10]:   Months      Brand      Total Sales ($)
0  09/2018    10x Infused      1,711.334232
1  09/2018  1964 Supply Co.    25,475.21594500000
2  09/2018    3 Bros Grow      120,153.644757
3  09/2018     3 Leaf    6,063.5297850000000
4  09/2018    350 Fire    631,510.0481550000
5  09/2018    710 Labs    2,065,970.9803990000
6  09/2018  A&A Craft Inc    5,094.305340000000
7  09/2018  AA Packaging    2,333.3399880000000
8  09/2018  Absolute Xtracts    5,747,227.563172000
9  09/2018   Aces Extracts    155,523.768684

```

```

[11]: brandTotalSales.describe()

```

```

[11]:   Months      Brand Total Sales ($)
count    25279      25279      25279
unique      37      1627      25277
top    05/2021  Island Cannabis Company      0
freq      848      37      3

```

```

[12]: brandTotalSales.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25279 entries, 0 to 25278
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Months      25279 non-null  object

```

```

1   Brand                25279 non-null  object
2   Total Sales ($)      25279 non-null  object
dtypes: object(3)
memory usage: 592.6+ KB

```

```
[13]: brandTotalUnits.head(10)
```

```

[13]:      Brands  Months  Total Units  vs. Prior Period
0   #BlackSeries  08/2020  1,616.3390040000000      NaN
1   #BlackSeries  09/2020      NaN      -1.000000
2   #BlackSeries  01/2021  715.5328380000000      NaN
3   #BlackSeries  02/2021    766.669135      0.071466
4   #BlackSeries  03/2021      NaN      -1.000000
5  101 Cannabis Co.  11/2019    131.06772      NaN
6  101 Cannabis Co.  12/2019      NaN      -1.000000
7  101 Cannabis Co.  01/2020  345.4134480000000      NaN
8  101 Cannabis Co.  02/2020  696.6584310000000      1.016883
9  101 Cannabis Co.  03/2020  943.3933280000000      0.354169

```

```
[14]: brandTotalUnits.describe()
```

```

[14]:      vs. Prior Period
count      24935.000000
mean         0.265306
std          3.291373
min         -1.000000
25%         -0.351822
50%         -0.055216
75%          0.240113
max         250.792020

```

```
[15]: brandTotalUnits.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27686 entries, 0 to 27685
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brands                27686 non-null  object
1   Months                27686 non-null  object
2   Total Units           25712 non-null  object
3   vs. Prior Period      24935 non-null  float64
dtypes: float64(1), object(3)
memory usage: 865.3+ KB

```

```

[16]: brands = brandTotalSales["Brand"].unique()
len(brands)

```


[16]: 1627

```
[17]: brandTotalSales['Total Sales ($)'] = brandTotalSales['Total Sales ($)'].str.  
      ↪replace(",", "")[0:]  
brandTotalSales['Total Sales ($)'] = pd.to_numeric(brandTotalSales['Total Sales_'  
      ↪($)'])  
  
brandTotalUnits['Total Units'] = brandTotalUnits['Total Units'].str.  
      ↪replace(",", "")[0:]  
brandTotalUnits['Total Units'] = pd.to_numeric(brandTotalUnits['Total Units'])  
  
brandDetails['Total THC'] = brandDetails['Total THC'].str.replace(",", "")[0:]  
brandDetails['Total THC'] = pd.to_numeric(brandDetails['Total THC'])  
  
brandDetails['Total CBD'] = brandDetails['Total CBD'].str.replace(",", "")[0:]  
brandDetails['Total CBD'] = pd.to_numeric(brandDetails['Total CBD'])  
  
brandTotalSales['Months'] = pd.to_datetime(brandTotalSales['Months'])  
brandTotalUnits['Months'] = pd.to_datetime(brandTotalUnits['Months'])  
brandAvgRetPrice['Months'] = pd.to_datetime(brandAvgRetPrice['Months'])  
  
brandTotalSales.head(10)  
  
filtered_data = pd.DataFrame()  
  
for brand in brands:  
  
    brand_sales = brandTotalSales[brandTotalSales.Brand == brand]  
    brand_sales.loc[:, 'Last Month'] = brand_sales.loc[:, 'Total Sales ($)'].  
    ↪shift(1)  
    brand_sales.loc[:, '3 Month Avg'] = (brand_sales.loc[:, 'Total Sales ($)'].  
    ↪shift(1) + brand_sales.loc[:, 'Total Sales ($)'].shift(2) + brand_sales.loc[:,  
    ↪, 'Total Sales ($)'].shift(3))/3  
    #brand_sales.loc[:, '12 Month Avg'] = (brand_sales.loc[:, 'Total Sales ($)'].  
    ↪shift(1) + brand_sales.loc[:, 'Total Sales ($)'].shift(2) + brand_sales.loc[:,  
    ↪, 'Total Sales ($)'].shift(3) + brand_sales.loc[:, 'Total Sales ($)'].shift(4) +  
    ↪+ brand_sales.loc[:, 'Total Sales ($)'].shift(5) + brand_sales.loc[:, 'Total_  
    ↪Sales ($)'].shift(6) + brand_sales.loc[:, 'Total Sales ($)'].shift(7) +  
    ↪brand_sales.loc[:, 'Total Sales ($)'].shift(8) + brand_sales.loc[:, 'Total_  
    ↪Sales ($)'].shift(9) + brand_sales.loc[:, 'Total Sales ($)'].shift(10) +  
    ↪brand_sales.loc[:, 'Total Sales ($)'].shift(11) + brand_sales.loc[:, 'Total_  
    ↪Sales ($)'].shift(12))/12  
    brand_sales.loc[:, 'Last Year'] = (brand_sales.loc[:, 'Total Sales ($)'].  
    ↪shift(12))  
  
    brand_units = brandTotalUnits[brandTotalUnits.Brands == brand]
```

```

merged_data = brand_sales.merge(brand_units, left_on='Months',
↪right_on='Months')
merged_data = merged_data.drop(['Brands'], 1)

arp_data = brandAvgRetPrice[brandAvgRetPrice.Brands == brand]
merged_data = merged_data.merge(arp_data, left_on='Months',
↪right_on='Months')
merged_data = merged_data.drop(['Brands'], 1)

merged_data = merged_data.drop(['vs. Prior Period_x'], 1)
merged_data = merged_data.drop(['vs. Prior Period_y'], 1)

merged_data['Month_Num'] = merged_data['Months'].dt.month

one_hot_encoded_months = pd.get_dummies(merged_data['Month_Num'])

merged_data = pd.concat([merged_data, one_hot_encoded_months], axis=1)
merged_data = merged_data.drop(['Month_Num'], 1)

# Feature Engineering

brand_details = brandDetails[brandDetails.Brand == brand]

if len(brand_details) == 0:
    merged_data['Pax_Filter'] = float('NaN')
    merged_data['Inhaleables'] = float('NaN')
    merged_data['Ingestibles'] = float('NaN')
    merged_data['Other_Cannabis'] = float('NaN')
    merged_data['Topicals'] = float('NaN')
    merged_data['Max_THC'] = float('NaN')
    merged_data['Max_CBD'] = float('NaN')
    merged_data['Sell_CBD'] = float('NaN')
    merged_data['Mean_IPP'] = float('NaN')
    merged_data['Product_Count'] = float('NaN')
    merged_data['Price_Range'] = float('NaN')
    merged_data['Flavored_Count'] = float('NaN')
    merged_data['Mood_Count'] = float('NaN')

    filtered_data = filtered_data.append(merged_data)
    continue

inhaleables_sold = 0

```

```

ingestibles_sold = 0
other_sold = 0
topicals_sold = 0

if 'Inhaleables' in brand_details['Category L1'].values:
    inhaleables_sold = 1
if 'Ingestibles' in brand_details['Category L1'].values:
    ingestibles_sold = 1
if 'Other Cannabis' in brand_details['Category L1'].values:
    other_sold = 1
if 'Topicals' in brand_details['Category L1'].values:
    topicals_sold = 1
if 'All accessories' in brand_details['Category L1'].values:
    inhaleables_sold = 1
    ingestibles_sold = 1
    other_sold = 1
    topicals_sold = 1

pax_filter = float('NaN')
if 'Pax' in brand_details['Pax Filter'].values:
    pax_filter = 1
elif 'Not Pax' in brand_details['Pax Filter'].values:
    pax_filter = 0

merged_data['Pax_Filter'] = pax_filter

merged_data['Inhaleables'] = inhaleables_sold
merged_data['Ingestibles'] = ingestibles_sold
merged_data['Other_Cannabis'] = other_sold
merged_data['Topicals'] = topicals_sold

max_product_thc = 0
max_product_cbd = 0

avg_product_thc = 0
avg_product_cbd = 0

if brand_details['Total THC'].count() != 0:
    max_product_thc = max(brand_details['Total THC'])
    avg_product_thc = statistics.mean(brand_details['Total THC'])

if brand_details['Total CBD'].count() != 0:
    max_product_cbd = max(brand_details['Total CBD'])
    avg_product_cbd = statistics.mean(brand_details['Total CBD'])

```

```

contains_cbd = 0

for x in brand_details['Contains CBD'].values:
    if x != 'THC Only':
        contains_cbd = 1

merged_data['Max_THC'] = max_product_thc
merged_data['Max_CBD'] = max_product_cbd
merged_data['Mean_THC'] = avg_product_thc
merged_data['Mean_CBD'] = avg_product_cbd

merged_data['Sell_CBD'] = contains_cbd

mean_ipp = 1
if brand_details['Items Per Pack'].count() != 0:
    mean_ipp = statistics.mean(brand_details['Items Per Pack'])

merged_data['Mean_IPP'] = mean_ipp

product_count = len(brand_details)

merged_data['Product_Count'] = product_count

price_range = len(brand_details['$5 Price Increment'].unique())
#price_range = len(price_range_set)
merged_data['Price_Range'] = price_range

flavored_count = 0
mood_count = 0

for x in brand_details['Is Flavored'].values:
    if x == 'Flavored':
        flavored_count += 1
for x in brand_details['Mood Effect'].values:
    if x == 'Mood Specific':
        mood_count += 1

merged_data['Flavored_Count'] = flavored_count
merged_data['Mood_Count'] = mood_count

filtered_data = filtered_data.append(merged_data)
#print(str(filtered_data.shape))

```

/Users/nbarron/opt/anaconda3/lib/python3.8/site-packages/pandas/core/indexing.py:1597: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.obj[key] = value
/Users/nbarron/opt/anaconda3/lib/python3.8/site-
packages/pandas/core/indexing.py:1676: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
[18]: filtered_data.head(50)
```

```
[18]:
```

	Months	Brand	Total Sales (\$)	Last Month	3 Month Avg	\
0	2018-09-01	10x Infused	1711.334232	NaN	NaN	
0	2018-09-01	1964 Supply Co.	25475.215945	NaN	NaN	
1	2018-10-01	1964 Supply Co.	13613.214128	25475.215945	NaN	
2	2018-11-01	1964 Supply Co.	5402.873064	13613.214128	NaN	
3	2018-12-01	1964 Supply Co.	11862.458357	5402.873064	14830.434379	
4	2019-01-01	1964 Supply Co.	3999.035205	11862.458357	10292.848516	
5	2019-02-01	1964 Supply Co.	2417.479974	3999.035205	7088.122209	
6	2019-03-01	1964 Supply Co.	1607.563310	2417.479974	6092.991179	
7	2019-04-01	1964 Supply Co.	292.135879	1607.563310	2674.692830	
0	2018-09-01	3 Bros Grow	120153.644757	NaN	NaN	
1	2018-10-01	3 Bros Grow	112932.164895	120153.644757	NaN	
2	2018-11-01	3 Bros Grow	109432.452831	112932.164895	NaN	
3	2018-12-01	3 Bros Grow	208424.645419	109432.452831	114172.754161	
4	2019-01-01	3 Bros Grow	214650.825825	208424.645419	143596.421048	
5	2019-02-01	3 Bros Grow	557059.818673	214650.825825	177502.641358	
6	2019-03-01	3 Bros Grow	346319.611428	557059.818673	326711.763306	
7	2019-04-01	3 Bros Grow	519579.324303	346319.611428	372676.751975	
8	2019-05-01	3 Bros Grow	278252.378245	519579.324303	474319.584801	
9	2019-06-01	3 Bros Grow	132809.898997	278252.378245	381383.771325	
10	2019-07-01	3 Bros Grow	95303.780457	132809.898997	310213.867182	
11	2019-08-01	3 Bros Grow	120435.066126	95303.780457	168788.685900	
12	2019-09-01	3 Bros Grow	444813.781709	120435.066126	116182.915193	
13	2019-10-01	3 Bros Grow	323920.510121	444813.781709	220184.209431	
14	2019-11-01	3 Bros Grow	163786.306082	323920.510121	296389.785985	
15	2019-12-01	3 Bros Grow	409535.828011	163786.306082	310840.199304	
16	2020-01-01	3 Bros Grow	466658.723817	409535.828011	299080.881405	
17	2020-02-01	3 Bros Grow	227941.631626	466658.723817	346660.285970	
18	2020-03-01	3 Bros Grow	111775.989595	227941.631626	368045.394485	
19	2020-04-01	3 Bros Grow	120002.708661	111775.989595	268792.115013	
20	2020-05-01	3 Bros Grow	211402.393646	120002.708661	153240.109961	
21	2020-06-01	3 Bros Grow	155015.603879	211402.393646	147727.030634	
22	2020-07-01	3 Bros Grow	38776.507573	155015.603879	162140.235395	

23	2020-08-01	3 Bros Grow	56271.446989	38776.507573	135064.835033
24	2020-09-01	3 Bros Grow	276099.258355	56271.446989	83354.519480
25	2020-10-01	3 Bros Grow	23070.759975	276099.258355	123715.737639
26	2020-11-01	3 Bros Grow	3073.929764	23070.759975	118480.488440
27	2020-12-01	3 Bros Grow	2446.050652	3073.929764	100747.982698
28	2021-01-01	3 Bros Grow	32326.493793	2446.050652	9530.246797
0	2018-09-01	3 Leaf	6063.529785	NaN	NaN
1	2018-10-01	3 Leaf	27349.643956	6063.529785	NaN
2	2018-11-01	3 Leaf	44414.218569	27349.643956	NaN
3	2018-12-01	3 Leaf	43533.767285	44414.218569	25942.464103
4	2019-01-01	3 Leaf	8971.384508	43533.767285	38432.543270
5	2019-02-01	3 Leaf	12853.649203	8971.384508	32306.456787
6	2019-03-01	3 Leaf	14397.597306	12853.649203	21786.266999
7	2019-04-01	3 Leaf	12897.159109	14397.597306	12074.210339
8	2019-05-01	3 Leaf	9865.215140	12897.159109	13382.801873
9	2019-06-01	3 Leaf	14415.335434	9865.215140	12386.657185
10	2019-07-01	3 Leaf	22075.753500	14415.335434	12392.569894
11	2019-08-01	3 Leaf	18285.969875	22075.753500	15452.101358

	Last Year	Total Units	ARP	9	Pax_Filter	...	4	10	\
0	NaN	142.839336	11.980833	1.0	NaN	...	NaN	NaN	
0	NaN	2395.534726	10.634459	1.0	NaN	...	0.0	0.0	
1	NaN	1910.329288	7.126109	0.0	NaN	...	0.0	1.0	
2	NaN	502.815600	10.745238	0.0	NaN	...	0.0	0.0	
3	NaN	2251.347983	5.269047	0.0	NaN	...	0.0	0.0	
4	NaN	379.006944	10.551351	0.0	NaN	...	0.0	0.0	
5	NaN	225.686193	10.711688	0.0	NaN	...	0.0	0.0	
6	NaN	179.313911	8.965079	0.0	NaN	...	0.0	0.0	
7	NaN	39.325986	7.428571	0.0	NaN	...	1.0	0.0	
0	NaN	10018.989140	11.992592	1.0	NaN	...	0.0	0.0	
1	NaN	11214.910342	10.069823	0.0	NaN	...	0.0	1.0	
2	NaN	10439.708895	10.482328	0.0	NaN	...	0.0	0.0	
3	NaN	24908.365990	8.367656	0.0	NaN	...	0.0	0.0	
4	NaN	20666.853248	10.386237	0.0	NaN	...	0.0	0.0	
5	NaN	54397.406606	10.240558	0.0	NaN	...	0.0	0.0	
6	NaN	25289.238229	13.694347	0.0	NaN	...	0.0	0.0	
7	NaN	45758.593710	11.354792	0.0	NaN	...	1.0	0.0	
8	NaN	22375.374776	12.435652	0.0	NaN	...	0.0	0.0	
9	NaN	10138.122486	13.100049	0.0	NaN	...	0.0	0.0	
10	NaN	9047.681531	10.533503	0.0	NaN	...	0.0	0.0	
11	NaN	9049.551241	13.308402	0.0	NaN	...	0.0	0.0	
12	120153.644757	37202.683417	11.956497	1.0	NaN	...	0.0	0.0	
13	112932.164895	26769.936733	12.100160	0.0	NaN	...	0.0	1.0	
14	109432.452831	10463.572980	15.653000	0.0	NaN	...	0.0	0.0	
15	208424.645419	33470.498126	12.235726	0.0	NaN	...	0.0	0.0	
16	214650.825825	24020.626967	19.427416	0.0	NaN	...	0.0	0.0	
17	557059.818673	15279.093831	14.918531	0.0	NaN	...	0.0	0.0	

18	346319.611428	9886.948888	11.305408	0.0	NaN	...	0.0	0.0
19	519579.324303	12702.573524	9.447118	0.0	NaN	...	1.0	0.0
20	278252.378245	18189.435752	11.622262	0.0	NaN	...	0.0	0.0
21	132809.898997	13677.707625	11.333449	0.0	NaN	...	0.0	0.0
22	95303.780457	2636.133742	14.709613	0.0	NaN	...	0.0	0.0
23	120435.066126	3821.810916	14.723765	0.0	NaN	...	0.0	0.0
24	444813.781709	26713.549147	10.335551	1.0	NaN	...	0.0	0.0
25	323920.510121	1789.963356	12.888957	0.0	NaN	...	0.0	1.0
26	163786.306082	67.751596	45.370588	0.0	NaN	...	0.0	0.0
27	409535.828011	267.651935	9.138924	0.0	NaN	...	0.0	0.0
28	466658.723817	2862.131300	11.294553	0.0	NaN	...	0.0	0.0
0	NaN	1101.053215	5.507027	1.0	NaN	...	0.0	0.0
1	NaN	4081.949816	6.700142	0.0	NaN	...	0.0	1.0
2	NaN	6809.559840	6.522333	0.0	NaN	...	0.0	0.0
3	NaN	6833.258003	6.370865	0.0	NaN	...	0.0	0.0
4	NaN	1669.386544	5.374061	0.0	NaN	...	0.0	0.0
5	NaN	2022.382755	6.355696	0.0	NaN	...	0.0	0.0
6	NaN	2755.172420	5.225661	0.0	NaN	...	0.0	0.0
7	NaN	3000.010932	4.299037	0.0	NaN	...	1.0	0.0
8	NaN	2474.059833	3.987460	0.0	NaN	...	0.0	0.0
9	NaN	2571.203232	5.606455	0.0	NaN	...	0.0	0.0
10	NaN	3879.294294	5.690662	0.0	NaN	...	0.0	0.0
11	NaN	3013.540740	6.067935	0.0	NaN	...	0.0	0.0

	11	12	5	6	7	8	Mean_THC	Mean_CBD
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
0	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
1	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
2	1.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
3	0.0	1.0	NaN	NaN	NaN	NaN	NaN	NaN
4	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
5	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
6	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
7	0.0	0.0	NaN	NaN	NaN	NaN	NaN	NaN
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
8	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
9	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
10	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
11	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
12	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

13	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
14	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
15	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
16	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
17	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
18	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
19	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
20	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
21	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
22	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
23	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
24	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
25	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
26	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
27	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
28	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
1	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
2	1.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
3	0.0	1.0	0.0	0.0	0.0	0.0	NaN	NaN
4	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
5	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
6	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
7	0.0	0.0	0.0	0.0	0.0	0.0	NaN	NaN
8	0.0	0.0	1.0	0.0	0.0	0.0	NaN	NaN
9	0.0	0.0	0.0	1.0	0.0	0.0	NaN	NaN
10	0.0	0.0	0.0	0.0	1.0	0.0	NaN	NaN
11	0.0	0.0	0.0	0.0	0.0	1.0	NaN	NaN

[50 rows x 35 columns]

```
[19]: print(filtered_data.columns.tolist())
```

```
['Months', 'Brand', 'Total Sales ($)', 'Last Month', '3 Month Avg', 'Last Year',
'Total Units', 'ARP', 9, 'Pax_Filter', 'Inhaleables', 'Ingestibles',
'Other_Cannabis', 'Topicals', 'Max_THC', 'Max_CBD', 'Sell_CBD', 'Mean_IPP',
'Product_Count', 'Price_Range', 'Flavored_Count', 'Mood_Count', 1, 2, 3, 4, 10,
11, 12, 5, 6, 7, 8, 'Mean_THC', 'Mean_CBD']
```

```
[20]: filtered_data['1'] = filtered_data[1].fillna(0)
filtered_data['2'] = filtered_data[2].fillna(0)
filtered_data['3'] = filtered_data[3].fillna(0)
filtered_data['4'] = filtered_data[4].fillna(0)
filtered_data['5'] = filtered_data[5].fillna(0)
filtered_data['6'] = filtered_data[6].fillna(0)
filtered_data['7'] = filtered_data[7].fillna(0)
filtered_data['8'] = filtered_data[8].fillna(0)
filtered_data['9'] = filtered_data[9].fillna(0)
```



```

filtered_data['10'] = filtered_data[10].fillna(0)
filtered_data['11'] = filtered_data[11].fillna(0)
filtered_data['12'] = filtered_data[12].fillna(0)

filtered_data = filtered_data.drop([1], 1)
filtered_data = filtered_data.drop([2], 1)
filtered_data = filtered_data.drop([3], 1)
filtered_data = filtered_data.drop([4], 1)
filtered_data = filtered_data.drop([5], 1)
filtered_data = filtered_data.drop([6], 1)
filtered_data = filtered_data.drop([7], 1)
filtered_data = filtered_data.drop([8], 1)
filtered_data = filtered_data.drop([9], 1)
filtered_data = filtered_data.drop([10], 1)
filtered_data = filtered_data.drop([11], 1)
filtered_data = filtered_data.drop([12], 1)

```

```
[21]: filtered_data.head(50)
```

```

[21]:
   Months      Brand  Total Sales ($)  Last Month  3 Month Avg \
0  2018-09-01    10x Infused      1711.334232         NaN         NaN
0  2018-09-01  1964 Supply Co.     25475.215945         NaN         NaN
1  2018-10-01  1964 Supply Co.     13613.214128    25475.215945         NaN
2  2018-11-01  1964 Supply Co.       5402.873064    13613.214128         NaN
3  2018-12-01  1964 Supply Co.     11862.458357     5402.873064    14830.434379
4  2019-01-01  1964 Supply Co.       3999.035205     11862.458357    10292.848516
5  2019-02-01  1964 Supply Co.       2417.479974       3999.035205     7088.122209
6  2019-03-01  1964 Supply Co.       1607.563310       2417.479974     6092.991179
7  2019-04-01  1964 Supply Co.        292.135879       1607.563310     2674.692830
0  2018-09-01     3 Bros Grow     120153.644757         NaN         NaN
1  2018-10-01     3 Bros Grow     112932.164895    120153.644757         NaN
2  2018-11-01     3 Bros Grow     109432.452831    112932.164895         NaN
3  2018-12-01     3 Bros Grow     208424.645419    109432.452831    114172.754161
4  2019-01-01     3 Bros Grow     214650.825825    208424.645419    143596.421048
5  2019-02-01     3 Bros Grow     557059.818673    214650.825825    177502.641358
6  2019-03-01     3 Bros Grow     346319.611428    557059.818673    326711.763306
7  2019-04-01     3 Bros Grow     519579.324303    346319.611428    372676.751975
8  2019-05-01     3 Bros Grow     278252.378245    519579.324303    474319.584801
9  2019-06-01     3 Bros Grow     132809.898997    278252.378245    381383.771325
10 2019-07-01     3 Bros Grow      95303.780457    132809.898997    310213.867182
11 2019-08-01     3 Bros Grow     120435.066126     95303.780457    168788.685900
12 2019-09-01     3 Bros Grow     444813.781709    120435.066126    116182.915193
13 2019-10-01     3 Bros Grow     323920.510121    444813.781709    220184.209431
14 2019-11-01     3 Bros Grow     163786.306082    323920.510121    296389.785985
15 2019-12-01     3 Bros Grow     409535.828011    163786.306082    310840.199304
16 2020-01-01     3 Bros Grow     466658.723817    409535.828011    299080.881405
17 2020-02-01     3 Bros Grow     227941.631626    466658.723817    346660.285970

```

18	2020-03-01	3 Bros Grow	111775.989595	227941.631626	368045.394485
19	2020-04-01	3 Bros Grow	120002.708661	111775.989595	268792.115013
20	2020-05-01	3 Bros Grow	211402.393646	120002.708661	153240.109961
21	2020-06-01	3 Bros Grow	155015.603879	211402.393646	147727.030634
22	2020-07-01	3 Bros Grow	38776.507573	155015.603879	162140.235395
23	2020-08-01	3 Bros Grow	56271.446989	38776.507573	135064.835033
24	2020-09-01	3 Bros Grow	276099.258355	56271.446989	83354.519480
25	2020-10-01	3 Bros Grow	23070.759975	276099.258355	123715.737639
26	2020-11-01	3 Bros Grow	3073.929764	23070.759975	118480.488440
27	2020-12-01	3 Bros Grow	2446.050652	3073.929764	100747.982698
28	2021-01-01	3 Bros Grow	32326.493793	2446.050652	9530.246797
0	2018-09-01	3 Leaf	6063.529785	NaN	NaN
1	2018-10-01	3 Leaf	27349.643956	6063.529785	NaN
2	2018-11-01	3 Leaf	44414.218569	27349.643956	NaN
3	2018-12-01	3 Leaf	43533.767285	44414.218569	25942.464103
4	2019-01-01	3 Leaf	8971.384508	43533.767285	38432.543270
5	2019-02-01	3 Leaf	12853.649203	8971.384508	32306.456787
6	2019-03-01	3 Leaf	14397.597306	12853.649203	21786.266999
7	2019-04-01	3 Leaf	12897.159109	14397.597306	12074.210339
8	2019-05-01	3 Leaf	9865.215140	12897.159109	13382.801873
9	2019-06-01	3 Leaf	14415.335434	9865.215140	12386.657185
10	2019-07-01	3 Leaf	22075.753500	14415.335434	12392.569894
11	2019-08-01	3 Leaf	18285.969875	22075.753500	15452.101358

	Last Year	Total Units	ARP	Pax_Filter	Inhaleables	...	3 \
0	NaN	142.839336	11.980833	NaN	NaN	...	0.0
0	NaN	2395.534726	10.634459	NaN	NaN	...	0.0
1	NaN	1910.329288	7.126109	NaN	NaN	...	0.0
2	NaN	502.815600	10.745238	NaN	NaN	...	0.0
3	NaN	2251.347983	5.269047	NaN	NaN	...	0.0
4	NaN	379.006944	10.551351	NaN	NaN	...	0.0
5	NaN	225.686193	10.711688	NaN	NaN	...	0.0
6	NaN	179.313911	8.965079	NaN	NaN	...	1.0
7	NaN	39.325986	7.428571	NaN	NaN	...	0.0
0	NaN	10018.989140	11.992592	NaN	1.0	...	0.0
1	NaN	11214.910342	10.069823	NaN	1.0	...	0.0
2	NaN	10439.708895	10.482328	NaN	1.0	...	0.0
3	NaN	24908.365990	8.367656	NaN	1.0	...	0.0
4	NaN	20666.853248	10.386237	NaN	1.0	...	0.0
5	NaN	54397.406606	10.240558	NaN	1.0	...	0.0
6	NaN	25289.238229	13.694347	NaN	1.0	...	1.0
7	NaN	45758.593710	11.354792	NaN	1.0	...	0.0
8	NaN	22375.374776	12.435652	NaN	1.0	...	0.0
9	NaN	10138.122486	13.100049	NaN	1.0	...	0.0
10	NaN	9047.681531	10.533503	NaN	1.0	...	0.0
11	NaN	9049.551241	13.308402	NaN	1.0	...	0.0
12	120153.644757	37202.683417	11.956497	NaN	1.0	...	0.0

13	112932.164895	26769.936733	12.100160	NaN	1.0	...	0.0
14	109432.452831	10463.572980	15.653000	NaN	1.0	...	0.0
15	208424.645419	33470.498126	12.235726	NaN	1.0	...	0.0
16	214650.825825	24020.626967	19.427416	NaN	1.0	...	0.0
17	557059.818673	15279.093831	14.918531	NaN	1.0	...	0.0
18	346319.611428	9886.948888	11.305408	NaN	1.0	...	1.0
19	519579.324303	12702.573524	9.447118	NaN	1.0	...	0.0
20	278252.378245	18189.435752	11.622262	NaN	1.0	...	0.0
21	132809.898997	13677.707625	11.333449	NaN	1.0	...	0.0
22	95303.780457	2636.133742	14.709613	NaN	1.0	...	0.0
23	120435.066126	3821.810916	14.723765	NaN	1.0	...	0.0
24	444813.781709	26713.549147	10.335551	NaN	1.0	...	0.0
25	323920.510121	1789.963356	12.888957	NaN	1.0	...	0.0
26	163786.306082	67.751596	45.370588	NaN	1.0	...	0.0
27	409535.828011	267.651935	9.138924	NaN	1.0	...	0.0
28	466658.723817	2862.131300	11.294553	NaN	1.0	...	0.0
0	NaN	1101.053215	5.507027	NaN	NaN	...	0.0
1	NaN	4081.949816	6.700142	NaN	NaN	...	0.0
2	NaN	6809.559840	6.522333	NaN	NaN	...	0.0
3	NaN	6833.258003	6.370865	NaN	NaN	...	0.0
4	NaN	1669.386544	5.374061	NaN	NaN	...	0.0
5	NaN	2022.382755	6.355696	NaN	NaN	...	0.0
6	NaN	2755.172420	5.225661	NaN	NaN	...	1.0
7	NaN	3000.010932	4.299037	NaN	NaN	...	0.0
8	NaN	2474.059833	3.987460	NaN	NaN	...	0.0
9	NaN	2571.203232	5.606455	NaN	NaN	...	0.0
10	NaN	3879.294294	5.690662	NaN	NaN	...	0.0
11	NaN	3013.540740	6.067935	NaN	NaN	...	0.0

	4	5	6	7	8	9	10	11	12
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

8	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
9	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
10	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
11	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
12	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
13	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
14	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
15	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
16	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
17	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
18	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
19	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
20	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
21	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
22	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
23	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
24	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
25	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
26	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
27	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
28	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
6	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
7	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
8	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
9	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
10	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
11	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0

[50 rows x 35 columns]

```
[22]: # Data Statistics

print("Average Monthly Sales: $" + str(statistics.mean(filtered_data['Total_
    ↳Sales ($)'])))
print("Monthly Sales Standard Deviation: $" + str(statistics.
    ↳stdev(filtered_data['Total Sales ($)'])))

print("Average Monthly Units Sold: " + str(statistics.mean(filtered_data['Total_
    ↳Units'])))
print("Monthly Units Sold Standard Deviation: " + str(statistics.
    ↳stdev(filtered_data['Total Units'])))
```

```
print("Average Monthly ARP: " + str(statistics.mean(filtered_data['ARP'])))
print("Monthly ARP Standard Deviation: " + str(statistics.
↳stdev(filtered_data['ARP'])))

filtered_data.describe()
```

Average Monthly Sales: \$409372.85619946336
Monthly Sales Standard Deviation: \$1596024.283035418
Average Monthly Units Sold: 28862.10067850273
Monthly Units Sold Standard Deviation: 161715.5821856867
Average Monthly ARP: 22.679731745813
Monthly ARP Standard Deviation: 19.802723938896023

```
[22]:
```

	Total Sales (\$)	Last Month	3 Month Avg	Last Year	\
count	2.527900e+04	2.365200e+04	2.073400e+04	1.142400e+04	
mean	4.093729e+05	4.245507e+05	4.551029e+05	5.516544e+05	
std	1.596024e+06	1.625582e+06	1.669072e+06	1.877350e+06	
min	0.000000e+00	0.000000e+00	6.011905e+01	0.000000e+00	
25%	1.390320e+04	1.608221e+04	2.249319e+04	3.295066e+04	
50%	6.210080e+04	6.905932e+04	8.316126e+04	1.227479e+05	
75%	2.473270e+05	2.627836e+05	2.920007e+05	3.909989e+05	
max	4.036351e+07	4.036351e+07	3.737876e+07	4.036351e+07	

	Total Units	ARP	Pax_Filter	Inhaleables	Ingestibles	\
count	2.527900e+04	25279.000000	5609.000000	21472.000000	21472.000000	
mean	2.886210e+04	22.679732	0.149224	0.705663	0.310404	
std	1.617156e+05	19.802724	0.356341	0.455755	0.462670	
min	3.842953e+00	0.000000	0.000000	0.000000	0.000000	
25%	7.169135e+02	10.512827	0.000000	0.000000	0.000000	
50%	3.605059e+03	17.033051	0.000000	1.000000	0.000000	
75%	1.564044e+04	31.505612	0.000000	1.000000	1.000000	
max	5.248082e+06	700.874984	1.000000	1.000000	1.000000	

	Other_Cannabis	...	3	4	5	\
count	21472.000000	...	25279.000000	25279.000000	25279.000000	
mean	0.059845	...	0.080660	0.082361	0.085802	
std	0.237206	...	0.272317	0.274919	0.280078	
min	0.000000	...	0.000000	0.000000	0.000000	
25%	0.000000	...	0.000000	0.000000	0.000000	
50%	0.000000	...	0.000000	0.000000	0.000000	
75%	0.000000	...	0.000000	0.000000	0.000000	
max	1.000000	...	1.000000	1.000000	1.000000	

	6	7	8	9	10	\
count	25279.000000	25279.000000	25279.000000	25279.000000	25279.000000	
mean	0.087147	0.086752	0.087701	0.103841	0.073856	

std	0.282057	0.281477	0.282866	0.305060	0.261541
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	11	12
count	25279.000000	25279.000000
mean	0.076981	0.077812
std	0.266566	0.267880
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

[8 rows x 33 columns]

```
[23]: half_one = filtered_data.iloc[:, [0,1,2,3,4,5,6,7,9,10,11,12,13,14,15,16]]
half_one.describe()
```

```
[23]:
```

	Total Sales (\$)	Last Month	3 Month Avg	Last Year \
count	2.527900e+04	2.365200e+04	2.073400e+04	1.142400e+04
mean	4.093729e+05	4.245507e+05	4.551029e+05	5.516544e+05
std	1.596024e+06	1.625582e+06	1.669072e+06	1.877350e+06
min	0.000000e+00	0.000000e+00	6.011905e+01	0.000000e+00
25%	1.390320e+04	1.608221e+04	2.249319e+04	3.295066e+04
50%	6.210080e+04	6.905932e+04	8.316126e+04	1.227479e+05
75%	2.473270e+05	2.627836e+05	2.920007e+05	3.909989e+05
max	4.036351e+07	4.036351e+07	3.737876e+07	4.036351e+07

	Total Units	ARP	Inhaleables	Ingestibles	Other_Cannabis \
count	2.527900e+04	25279.000000	21472.000000	21472.000000	21472.000000
mean	2.886210e+04	22.679732	0.705663	0.310404	0.059845
std	1.617156e+05	19.802724	0.455755	0.462670	0.237206
min	3.842953e+00	0.000000	0.000000	0.000000	0.000000
25%	7.169135e+02	10.512827	0.000000	0.000000	0.000000
50%	3.605059e+03	17.033051	1.000000	0.000000	0.000000
75%	1.564044e+04	31.505612	1.000000	1.000000	0.000000
max	5.248082e+06	700.874984	1.000000	1.000000	1.000000

	Topicals	Max_THC	Max_CBD	Sell_CBD	Mean_IPP
count	21472.000000	21472.000000	21472.000000	21472.000000	21472.000000
mean	0.093564	110.025079	99.128748	0.379238	2.697393
std	0.291228	287.258462	483.069571	0.485209	8.084836
min	0.000000	0.000000	0.000000	0.000000	0.000000

25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.253623
75%	0.000000	100.000000	0.000000	1.000000	2.130435
max	1.000000	2300.000000	10000.000000	1.000000	198.931250

```
[24]: half_two = filtered_data.iloc[:
↪, [17,18,19,20,21,22,23,24,25,26,8,27,28,29,30,31,32]]
half_two.describe()
```

```
[24]:
```

	Product_Count	Price_Range	Flavored_Count	Mood_Count	\
count	21472.000000	21472.000000	21472.000000	21472.000000	
mean	210.689316	10.781762	10.832340	10.568927	
std	516.966345	5.648736	39.652441	48.626711	
min	1.000000	1.000000	0.000000	0.000000	
25%	21.000000	6.000000	0.000000	0.000000	
50%	68.000000	10.000000	0.000000	0.000000	
75%	188.000000	15.000000	0.000000	0.000000	
max	9004.000000	22.000000	640.000000	672.000000	

	Mean_THC	Mean_CBD	1	2	3 \
count	21472.000000	21472.000000	25279.000000	25279.000000	25279.000000
mean	29.262824	13.666658	0.077495	0.079592	0.080660
std	74.689194	43.702023	0.267381	0.270665	0.272317
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	19.977802	0.000000	0.000000	0.000000	0.000000
max	1000.000000	425.581395	1.000000	1.000000	1.000000

	4	Pax_Filter	5	6	7 \
count	25279.000000	5609.000000	25279.000000	25279.000000	25279.000000
mean	0.082361	0.149224	0.085802	0.087147	0.086752
std	0.274919	0.356341	0.280078	0.282057	0.281477
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	8	9	10
count	25279.000000	25279.000000	25279.000000
mean	0.087701	0.103841	0.073856
std	0.282866	0.305060	0.261541
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000

max	1.000000	1.000000	1.000000
-----	----------	----------	----------

```
[25]: filtered_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25279 entries, 0 to 0
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Months                25279 non-null  datetime64[ns]
1   Brand                 25279 non-null  object
2   Total Sales ($)       25279 non-null  float64
3   Last Month            23652 non-null  float64
4   3 Month Avg           20734 non-null  float64
5   Last Year             11424 non-null  float64
6   Total Units           25279 non-null  float64
7   ARP                   25279 non-null  float64
8   Pax_Filter            5609 non-null   float64
9   Inhaleables           21472 non-null  float64
10  Ingestibles           21472 non-null  float64
11  Other_Cannabis        21472 non-null  float64
12  Topicals              21472 non-null  float64
13  Max_THC               21472 non-null  float64
14  Max_CBD               21472 non-null  float64
15  Sell_CBD              21472 non-null  float64
16  Mean_IPP              21472 non-null  float64
17  Product_Count         21472 non-null  float64
18  Price_Range           21472 non-null  float64
19  Flavored_Count        21472 non-null  float64
20  Mood_Count            21472 non-null  float64
21  Mean_THC              21472 non-null  float64
22  Mean_CBD              21472 non-null  float64
23  1                     25279 non-null  float64
24  2                     25279 non-null  float64
25  3                     25279 non-null  float64
26  4                     25279 non-null  float64
27  5                     25279 non-null  float64
28  6                     25279 non-null  float64
29  7                     25279 non-null  float64
30  8                     25279 non-null  float64
31  9                     25279 non-null  float64
32  10                    25279 non-null  float64
33  11                    25279 non-null  float64
34  12                    25279 non-null  float64
dtypes: datetime64[ns](1), float64(33), object(1)
memory usage: 6.9+ MB
```



```
[26]: # Feature Histograms

print("Total Monthly Sales ($) Histogram")
filtered_data['Total Sales ($)'].hist()
plt.yscale('log')
plt.show()

print("Total Monthly Units Histogram")
filtered_data['Total Units'].hist()
plt.yscale('log')
plt.show()

print("ARP Histogram")
filtered_data['ARP'].hist()
plt.yscale('log')
plt.show()

print("Pax Filter Histogram")
filtered_data['Pax_Filter'].hist()
plt.yscale('linear')
plt.show()

print("Inhaleables Sold Histogram")
filtered_data['Inhaleables'].hist()
plt.yscale('linear')
plt.show()

print("Ingenstibles Sold Histogram")
filtered_data['Ingestibles'].hist()
plt.yscale('linear')
plt.show()

print("Other Cannabis Product Sold Histogram")
filtered_data['Other_Cannabis'].hist()
plt.yscale('linear')
plt.show()

print("Topicals Sold Histogram")
filtered_data['Topicals'].hist()
plt.yscale('linear')
plt.show()

print("Max mg THC Histogram")
filtered_data['Max_THC'].hist()
plt.yscale('log')
plt.show()
```

```

print("Max mg CBD Histogram")
filtered_data['Max_CBD'].hist()
plt.yscale('log')
plt.show()

print("Product Count Histogram")
filtered_data['Product_Count'].hist()
plt.yscale('log')
plt.show()


print("Price Range Histogram")
filtered_data['Price_Range'].hist()
plt.yscale('linear')
plt.show()

print("Flavored Count Histogram")
filtered_data['Flavored_Count'].hist()
plt.yscale('log')
plt.show()


print("Mood Count Histogram")
filtered_data['Mood_Count'].hist()
plt.yscale('log')
plt.show()

```


Total Monthly Sales (\$) Histogram




Total Monthly Units Histogram




ARP Histogram




Pax Filter Histogram




Inhaleables Sold Histogram




Ingenstibles Sold Histogram




Other Cannabis Product Sold Histogram




Topicals Sold Histogram




Max mg THC Histogram




Max mg CBD Histogram




Product Count Histogram




Price Range Histogram



Flavored Count Histogram




Mood Count Histogram



```
[27]: plt.subplots(figsize=(20,15))  
sns.heatmap(filtered_data.corr(), cmap="PiYG")
```

```
[27]: <AxesSubplot:>
```



Findings

I originally suspected that certain months (numbered 1-12) would have a correlation with total sales (e.g. an average of higher sales in some months and lower in others), but the heatmap above suggests that this is not true.

Variables that have a strong correlation with sales include the following:

- Last Month Sales (positive correlation)
- Last 3 Month Average (positive Correlation)
- Last Year Sales (positive Correlation)
- Total Units Sold (positive correlation)
- Product Count (How many products are offered)
- Pax Filter Products Sold (weak positive correlation)
- Mean Product THC (weak negative correlation)
- Price Range (The range of price in products sold - the number different \$5 increment products available)
- Offering CBD products (weak positive correlation)
- CBD Product Offered (weak positive correlation)
- ARP (weak negative correlation)

Aside from the months, most other variables (Max CBD offered, etc) have either no correlation or a weak negative one with total sales

1 4. Additional Data Feature Extraction

Drop the following variables - month variables (1-12) - Mean_IPP - Mean_CBD - Max_CBD - Months - Brand - All other variables not present in altered dataset that are in original datasets - labels (units sold, total sales, total units)

Maintain Scalar values as is, binary values are already set. Category L1 has effectively been One Hot Encoded during preprocessing (not exactly, as more than one is possible, but the categorical values have been encoded into binary integer ones)

Imputation Strategies For Columns with NaN values - Last Month Set to 0 (no sales last month, new brand) - new column feature made for new brands - 3 Month Avg. Set to 0 - Last Year Set to 0 - Pax_Filter Set to 0

- Inhaleables Set to 0
- Ingestibles Set to 0
- Other_Cannabis: Set to 0
- Topicals: Set to 0
- Max_THC Set to Median
- Sell_CBD Set to 0
- Product_Count Set to Median
- Price_Range Set to Median
- Flavored_Count Set to Median
- Mood_Count Set to Median
- Mean_THC Set to Mean

Create new/cross features - New Product (Binary 1 or 0, depending on if Last Month is NaN or not) - Max THC and Ingestibles (These variables are positively correlated)

Scaling strategies:

Apply Standard Scaling to the following variables - Last Month - 3 Month Avg. - Last Year - Max_THC - Product_Count - Price_Range - Flavored_Count - Mood Count - Mean THC

Pipeline for all of these alterations is implemented directly below

```
[28]: # 5. Create Pipeline

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

data_y = filtered_data["Total Sales ($)"].copy()

med_impute_cols = ["Last Year", "Max_THC", "ARP", "Product_Count",
                  ↪ "Price_Range", "Flavored_Count", "Mood_Count", "MTHC_Ing_Cross"]
```

```

zero_scale_cols = []
zero_noscale_cols = ["Pax_Filter", "Inhaleables", "Ingestibles",
    ↳ "Other_Cannabis", "Topicals", "Sell_CBD", "New_Product"]
mean_impute_cols = ["Last Month", "3 Month Avg", "Mean_THC"]
#scale_cols = ["Last Month", "3 Month Avg", "Last Year", "Max_THC",
    ↳ "Product_Count", "Price_Range", "Flavored_Count", "Mood_Count", "Mean_THC",
    ↳ "MTHC_Ing_Cross"]

class AugmentFeatures(BaseEstimator, TransformerMixin):
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        X = X.drop(["1", "2", "3", "4", "5", "6", "7", "8", "9", "10", "11",
    ↳ "12", "Mean_IPP", "Mean_CBD", "Max_CBD", "Brand", "Total Sales ($)", "Total_
    ↳ Units", "Months"], 1)
        X['MTHC_Ing_Cross'] = X["Max_THC"] * X["Ingestibles"]
        X['New_Product'] = (float('-inf') < X["Last Month"]) & (float('inf') >
    ↳ X["Last Month"])
        X['New_Product'] = X['New_Product'].astype(int) * 1.0
        #X.info()
        return X

attr_filter = AugmentFeatures()
filtered_data_proc = attr_filter.transform(filtered_data)

med_pipeline = Pipeline([
    ('median_imputer', SimpleImputer(strategy='median')),
    ('std_scale', StandardScaler())
])

zero_scale_pipeline = Pipeline([
    ('zero_imputer', SimpleImputer(strategy='constant', fill_value=0)),
    ('std_scale', StandardScaler())
])

zero_noscale_pipeline = Pipeline([
    ('zero_imputer', SimpleImputer(strategy='constant', fill_value=0))
])

mean_pipeline = Pipeline([
    ('mean_imputer', SimpleImputer(strategy='mean')),
    ('std_scale', StandardScaler())
])

transformer = ColumnTransformer([
    ('med_imp', med_pipeline, med_impute_cols),
    ('zero_scale', zero_scale_pipeline, zero_scale_cols),

```

```

        ('zero_noscale', zero_noscale_pipeline, zero_noscale_cols),
        ('mean_imp', mean_pipeline, mean_impute_cols)
    ])

```

```

data_prepared = transformer.fit_transform(filtered_data_proc)

```

```

[29]: data_prepared.shape

```

```

[29]: (25279, 18)

```

```

[30]: # 7. Linear Regression

```

```

from sklearn.linear_model import LinearRegression
import sklearn.metrics as metrics
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV

X_train, X_test, y_train, y_test = train_test_split(data_prepared, data_y, \
    train_size=0.85, random_state=121)

lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

preds_train = lin_reg.predict(X_train)
preds_test = lin_reg.predict(X_test)

def regression_results(y_true, y_pred):
    # Regression metrics
    explained_variance=metrics.explained_variance_score(y_true, y_pred)
    mean_absolute_error=metrics.mean_absolute_error(y_true, y_pred)
    mse=metrics.mean_squared_error(y_true, y_pred)
    #mean_squared_log_error=metrics.mean_squared_log_error(y_true, y_pred)
    median_absolute_error=metrics.median_absolute_error(y_true, y_pred)
    r2=metrics.r2_score(y_true, y_pred)
    print('explained_variance: ', round(explained_variance,4))
    #print('mean_squared_log_error: ', round(mean_squared_log_error,4))
    print('r2: ', round(r2,4))
    print('MAE: ', round(mean_absolute_error,4))
    print('MSE: ', round(mse,4))
    print('RMSE: ', round(np.sqrt(mse),4))

print("Training Set Results: ")
regression_results(y_train, preds_train)

print('\n' + "Testing Set Results: ")
regression_results(y_test, preds_test)

```

Training Set Results:

```
explained_variance: 0.946
r2: 0.946
MAE: 82824.3978
MSE: 131439722934.8976
RMSE: 362546.1666
```

Testing Set Results:

```
explained_variance: 0.9776
r2: 0.9776
MAE: 89337.4836
MSE: 71251360719.1529
RMSE: 266929.5051
```

```
[31]: # Feature Importance based on weight in model
```

```
coefficients = lin_reg.coef_


feature_ordered_list = ["Last_Month", "3_Month_Avg", "Last_Year", "ARP",
    ↳ "Pax_Filer", "Inhaleables", "Ingestibles", "Other_Cannabis", "Topicals",
    ↳ "Max_THC", "Sell_CBD", "Product_Count", "Price_Range", "Flavored_Count",
    ↳ "Mood_Count", "Mean_THC", "MTHC_Ing_Cross", "New_Product"]


for i,v in enumerate(coefficients):
    print('Feature: %0d, %s Score: %.5f' % (i,feature_ordered_list[i],v))
# plot feature importance
plt.bar([x for x in range(len(coefficients))], coefficients)
plt.yscale('log')
plt.xlabel("feature number")
plt.ylabel("coefficient")
plt.show()

plt.bar([x for x in range(len(coefficients))], coefficients)
plt.yscale('linear')
plt.xlabel("feature number")
plt.ylabel("coefficient")
plt.show()
```

```
Feature: 0, Last_Month Score: -16726.93808
Feature: 1, 3_Month_Avg Score: -1252.40326
Feature: 2, Last_Year Score: -9613.25731
Feature: 3, ARP Score: 41837.51016
Feature: 4, Pax_Filer Score: 10005.31014
Feature: 5, Inhaleables Score: 2039.77175
Feature: 6, Ingestibles Score: 966.01786
Feature: 7, Other_Cannabis Score: -4060.54556
Feature: 8, Topicals Score: -6173.71553
Feature: 9, Max_THC Score: 13567.07529
Feature: 10, Sell_CBD Score: 12725.10978
```

Feature: 11, Product_Count Score: 27512.33976
Feature: 12, Price_Range Score: -10621.30227
Feature: 13, Flavored_Count Score: 18818.28732
Feature: 14, Mood_Count Score: 245711.08046
Feature: 15, Mean_THC Score: 1453753.73643
Feature: 16, MTHC_Ing_Cross Score: 81676.23301
Feature: 17, New_Product Score: -149.58918





```
[32]: import statsmodels.api as sm
from scipy import stats

X_train_2 = sm.add_constant(X_train)
est = sm.OLS(y_train, X_train_2)
lin_reg_2 = est.fit()
print(lin_reg_2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:          Total Sales ($)      R-squared:                0.946
Model:                  OLS                 Adj. R-squared:           0.946
Method:                 Least Squares       F-statistic:             2.089e+04
Date:                  Sat, 04 Dec 2021     Prob (F-statistic):       0.00
Time:                  01:28:48             Log-Likelihood:          -3.0554e+05
No. Observations:      21487                AIC:                    6.111e+05
Df Residuals:          21468                BIC:                    6.113e+05
Df Model:              18
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.616e+05	1.07e+04	15.154	0.000	1.41e+05	1.82e+05
x1	-1.673e+04	4277.383	-3.911	0.000	-2.51e+04	-8342.948

x2	-1252.4033	1.18e+04	-0.106	0.916	-2.45e+04	2.2e+04
x3	-9613.2573	2737.378	-3.512	0.000	-1.5e+04	-4247.792
x4	4.184e+04	3331.598	12.558	0.000	3.53e+04	4.84e+04
x5	1.001e+04	3166.037	3.160	0.002	3799.642	1.62e+04
x6	2039.7717	2988.370	0.683	0.495	-3817.656	7897.200
x7	966.0179	2587.289	0.373	0.709	-4105.261	6037.296
x8	-4060.5456	1.11e+04	-0.366	0.714	-2.58e+04	1.77e+04
x9	-6173.7155	1.53e+04	-0.403	0.687	-3.62e+04	2.38e+04
x10	1.357e+04	5931.178	2.287	0.022	1941.525	2.52e+04
x11	1.273e+04	8093.040	1.572	0.116	-3137.851	2.86e+04
x12	2.751e+04	1.2e+04	2.293	0.022	3995.970	5.1e+04
x13	-1.062e+04	1.15e+04	-0.923	0.356	-3.32e+04	1.19e+04
x14	1.882e+04	7134.866	2.638	0.008	4833.418	3.28e+04
x15	2.457e+05	1.03e+04	23.923	0.000	2.26e+05	2.66e+05
x16	1.454e+06	8437.734	172.292	0.000	1.44e+06	1.47e+06
x17	8.168e+04	9155.607	8.921	0.000	6.37e+04	9.96e+04
x18	-149.5892	4117.140	-0.036	0.971	-8219.491	7920.312

Omnibus:	75965.057	Durbin-Watson:	1.998
Prob(Omnibus):	0.000	Jarque-Bera (JB):	46929124358.997
Skew:	69.704	Prob(JB):	0.00
Kurtosis:	7241.664	Cond. No.	12.9

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[33]: # 8. PCA Analysis

from sklearn import decomposition
from sklearn.linear_model import LogisticRegression

r2_list = []
mae_list = []
mse_list = []
rmse_list = []

for i in range(1,18):

    pca = decomposition.PCA(n_components = i)
    data_pca = pca.fit_transform(data_prepared)
    X_train_pca, X_test_pca, y_train_pca, y_test_pca = \
    train_test_split(data_pca, data_y, train_size = 0.8)
    lin_reg_pca = LinearRegression()
    lin_reg_pca.fit(X_train_pca, y_train_pca)
    linreg_preds = lin_reg_pca.predict(X_test_pca)
```

```

r2 = metrics.r2_score(y_test_pca, linreg_preds)
mae = metrics.mean_absolute_error(y_test_pca, linreg_preds)
mse = metrics.mean_squared_error(y_test_pca, linreg_preds)
rmse = round(np.sqrt(mse),4)

r2_list.append(r2)
mae_list.append(mae)
mse_list.append(mse)
rmse_list.append(rmse)



plt.bar([x for x in range(len(r2_list))], r2_list)
plt.yscale('linear')
plt.xlabel("components")
plt.ylabel("R2")
plt.show()



plt.bar([x for x in range(len(mae_list))], mae_list)
plt.yscale('linear')
plt.xlabel("components")
plt.ylabel("MAE")
plt.show()

plt.bar([x for x in range(len(mse_list))], mse_list)
plt.yscale('linear')
plt.xlabel("components")
plt.ylabel("MSE")
plt.show()

plt.bar([x for x in range(len(rmse_list))], rmse_list)
plt.yscale('linear')
plt.xlabel("components")
plt.ylabel("RMSE")
plt.show()

```





Based on these results, I choose to keep 12 components in my dataset following PCA. Adding

additional components to this does not improve the performance of the model (looking at the data above tells us that increasing from 8 components does not really improve the model that much, but 12 is the closes number of components with improved performance in the model while still helping prevent overfitting).

```
[34]: pca = decomposition.PCA(n_components = 12)
data_pca = pca.fit_transform(data_prepared)
X_train_pca, X_test_pca, y_train_pca, y_test_pca = train_test_split(data_pca,
    ↪data_y, train_size = 0.8)

data_pca.shape
```

```
[34]: (25279, 12)
```

```
[35]: # 9. Ensemble Method
# Implement a KNN method for linear regression to have another baseline model,
    ↪to compare to

from sklearn.neighbors import KNeighborsRegressor

#neighbors = [10,50,100,150,200,300,500,750,1000,1500, 2000, 3000, 5000]
neighbors = [1,2,3,4,5,6,7,8,9,10]

r2_train_list = []
r2_test_list = []
mae_train_list = []
mae_test_list = []
mse_train_list = []
mse_test_list = []
rmse_train_list = []
rmse_test_list = []

for n in neighbors:
    knn_cur = KNeighborsRegressor(n_neighbors=n)
    knn_cur.fit(X_train_pca, y_train_pca)
    preds_train_cur = knn_cur.predict(X_train_pca)
    preds_test_cur = knn_cur.predict(X_test_pca)

    r2 = metrics.r2_score(y_train_pca, preds_train_cur)
    mae = metrics.mean_absolute_error(y_train_pca, preds_train_cur)
    mse = metrics.mean_squared_error(y_train_pca, preds_train_cur)
    rmse = round(np.sqrt(mse),4)

    r2_train_list.append(r2)
    mae_train_list.append(mae)
    mse_train_list.append(mse)
    rmse_train_list.append(rmse)
```

```

r2 = metrics.r2_score(y_test_pca, preds_test_cur)
mae = metrics.mean_absolute_error(y_test_pca, preds_test_cur)
mse = metrics.mean_squared_error(y_test_pca, preds_test_cur)
rmse = round(np.sqrt(mse),4)

r2_test_list.append(r2)
mae_test_list.append(mae)
mse_test_list.append(mse)
rmse_test_list.append(rmse)


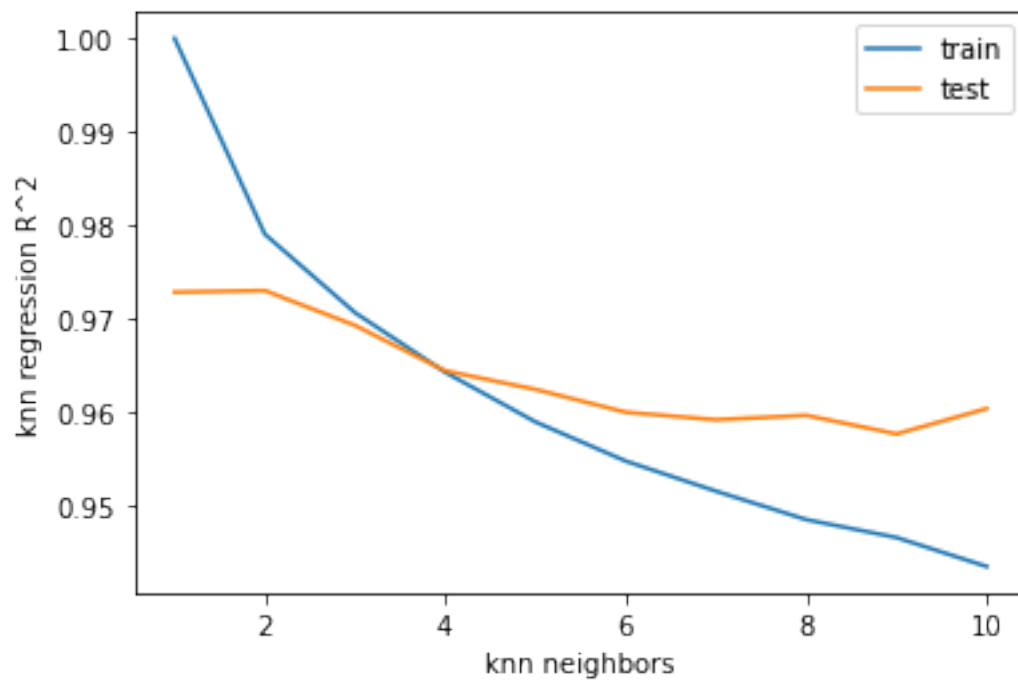
plt.xlabel('knn neighbors')
plt.ylabel('knn regression R^2')
plt.plot(neighbors, r2_train_list, label='train')
plt.plot(neighbors, r2_test_list, label='test')
plt.legend()
plt.show()



plt.xlabel('knn neighbors')
plt.ylabel('knn regression MAE')
plt.plot(neighbors, mae_train_list, label='train')
plt.plot(neighbors, mae_test_list, label='test')
plt.legend()
plt.show()

plt.xlabel('knn neighbors')
plt.ylabel('knn regression MSE')
plt.plot(neighbors, mse_train_list, label='train')
plt.plot(neighbors, mse_test_list, label='test')
plt.legend()
plt.show()

plt.xlabel('knn neighbors')
plt.ylabel('knn regression RMSE')
plt.plot(neighbors, rmse_train_list, label='train')
plt.plot(neighbors, rmse_test_list, label='test')
plt.legend()
plt.show()

```





Based on checking performance accuracy with multiple numbers of neighbors considered, using either 1 or 2 yields the greatest performance. We use 2 from now on.


```
[36]: knn_cur = KNeighborsRegressor(n_neighbors=2)
knn_cur.fit(X_train_pca, y_train_pca)
preds_train = knn_cur.predict(X_train_pca)
preds_test = knn_cur.predict(X_test_pca)

print("KNN Regression Model (n_neighbors=2) Metric Following PCA reduction to 12 components")

print("\nTraining Set Metrics")
regression_results(y_train_pca, preds_train)
print("\nTesting Set Metrics")
regression_results(y_test_pca, preds_test)

linreg_pca = LinearRegression()
linreg_pca.fit(X_train_pca, y_train_pca)
lin_reg_pca.fit(X_train_pca, y_train_pca)
linreg_preds_train = lin_reg_pca.predict(X_train_pca)
linreg_preds_test = lin_reg_pca.predict(X_test_pca)

print("Linear Regression Model Metric Following PCA reduction to 12 components")

print("\nTraining Set Metrics")
regression_results(y_train_pca, linreg_preds_train)
print("\nTesting Set Metrics")
regression_results(y_test_pca, linreg_preds_test)
```

KNN Regression Model (n_neighbors=2) Metric Following PCA reduction to 12 components

Training Set Metrics

explained_variance: 0.9791
r2: 0.9791
MAE: 49494.1261
MSE: 50756344592.9696
RMSE: 225291.6878

Testing Set Metrics

explained_variance: 0.9731
r2: 0.973
MAE: 84892.3058
MSE: 81707341523.4399
RMSE: 285844.9606

Linear Regression Model Metric Following PCA reduction to 12 components

Training Set Metrics

explained_variance: 0.9298
r2: 0.9298
MAE: 131665.8164

```
MSE: 170356210799.2939
RMSE: 412742.3056
```

```
Testing Set Metrics
explained_variance: 0.9631
r2: 0.9631
MAE: 128519.7399
MSE: 111738089097.2553
RMSE: 334272.4773
```

By comparing the testing set metrics of the KNN regression and Linear REgression, we see that Linear Regression displays better performance (greater R^2 , lower MAE, MSE, and RMSE)

```
[37]: # Random Forest Ensemble Method

from sklearn.ensemble import RandomForestRegressor

rforest_reg = RandomForestRegressor()
rforest_reg.fit(X_train_pca, y_train_pca)
rf_preds_train = rforest_reg.predict(X_train_pca)
rf_preds_test = rforest_reg.predict(X_test_pca)

print("Random Forest Regressor Metrics Following PCA reduction to 12_
      ↪components")

print("\nTraining Set Metrics")
regression_results(y_train_pca, linreg_preds_train)
print("\nTesting Set Metrics")
regression_results(y_test_pca, linreg_preds_test)
```

Random Forest Regressor Metrics Following PCA reduction to 12 components

```
Training Set Metrics
explained_variance: 0.9298
r2: 0.9298
MAE: 131665.8164
MSE: 170356210799.2939
RMSE: 412742.3056
```

```
Testing Set Metrics
explained_variance: 0.9631
r2: 0.9631
MAE: 128519.7399
MSE: 111738089097.2553
RMSE: 334272.4773
```

```
[38]: # 10. Cross Validation (10-fold)
```

```

from sklearn.model_selection import KFold

kfold = KFold(n_splits = 10, shuffle=True)
splits = kfold.split(data_pca, data_y)

linreg_r2 = []
linreg_mae = []
linreg_mse = []
linreg_rmse = []

knnreg_r2 = []
knnreg_mae = []
knnreg_mse = []
knnreg_rmse = []

randfor_r2 = []
randfor_mae = []
randfor_mse = []
randfor_rmse = []

for train_indeces, test_indeces in splits:

    X_train = data_pca[train_indeces]
    y_train = data_y.iloc[train_indeces]

    X_test = data_pca[test_indeces, :]
    y_test = data_y.iloc[test_indeces]

    linreg = LinearRegression()
    knnreg = KNeighborsRegressor(n_neighbors=2)
    rafreg = RandomForestRegressor()

    linreg.fit(X_train, y_train)
    knnreg.fit(X_train, y_train)
    rafreg.fit(X_train, y_train)

    linreg_preds = linreg.predict(X_test)
    knnreg_preds = knnreg.predict(X_test)
    rafreg_preds = rafreg.predict(X_test)

    linreg_r2.append(metrics.r2_score(y_test, linreg_preds))
    linreg_mae.append(metrics.mean_absolute_error(y_test, linreg_preds))
    linreg_mse.append(metrics.mean_squared_error(y_test, linreg_preds))
    linreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪linreg_preds)),4))

```

```

knnreg_r2.append(metrics.r2_score(y_test, knnreg_preds))
knnreg_mae.append(metrics.mean_absolute_error(y_test, knnreg_preds))
knnreg_mse.append(metrics.mean_squared_error(y_test, knnreg_preds))
knnreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
→knnreg_preds)),4))

randfor_r2.append(metrics.r2_score(y_test, rafreg_preds))
randfor_mae.append(metrics.mean_absolute_error(y_test, rafreg_preds))
randfor_mse.append(metrics.mean_squared_error(y_test, rafreg_preds))
randfor_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
→rafreg_preds)),4))

print("Linear Regression Average R^2: " + str(np.mean(linreg_r2)))
print("Linear Regression Average MAE: " + str(np.mean(linreg_mae)))
print("Linear Regression Average MSE: " + str(np.mean(linreg_mse)))
print("Linear Regression Average RMSE: " + str(np.mean(linreg_rmse)))

print("KNN Regression Average R^2: " + str(np.mean(knnreg_r2)))
print("KNN Regression Average MAE: " + str(np.mean(knnreg_mae)))
print("KNN Regression Average MSE: " + str(np.mean(knnreg_mse)))
print("KNN Regression Average RMSE: " + str(np.mean(knnreg_rmse)))

print("Random Forest Regression Average R^2: " + str(np.mean(randfor_r2)))
print("Random Forest Average MAE: " + str(np.mean(randfor_mae)))
print("Random Forest Average MSE: " + str(np.mean(randfor_mse)))
print("Random Forest Average RMSE: " + str(np.mean(randfor_rmse)))

```

```

Linear Regression Average R^2: 0.9397111976739853
Linear Regression Average MAE: 131890.71095575058
Linear Regression Average MSE: 160319480077.21356
Linear Regression Average RMSE: 363027.2107
KNN Regression Average R^2: 0.9521348314837568
KNN Regression Average MAE: 85917.5833494429
KNN Regression Average MSE: 128593020773.09123
KNN Regression Average RMSE: 319363.48352
Random Forest Regression Average R^2: 0.948409285946294
Random Forest Average MAE: 87530.34380017841
Random Forest Average MSE: 137768222451.0931
Random Forest Average RMSE: 331004.81853000005

```

Compare this to data without PCA applied

```

[39]: from sklearn.model_selection import KFold

kfold = KFold(n_splits = 10, shuffle=True)
splits = kfold.split(data_prepared, data_y)

```

```

linreg_r2 = []
linreg_mae = []
linreg_mse = []
linreg_rmse = []

knnreg_r2 = []
knnreg_mae = []
knnreg_mse = []
knnreg_rmse = []

randfor_r2 = []
randfor_mae = []
randfor_mse = []
randfor_rmse = []

for train_indices, test_indices in splits:

    X_train = data_prepared[train_indices]
    y_train = data_y.iloc[train_indices]

    X_test = data_prepared[test_indices, :]
    y_test = data_y.iloc[test_indices]

    linreg = LinearRegression()
    knnreg = KNeighborsRegressor(n_neighbors=2)
    rafreg = RandomForestRegressor()

    linreg.fit(X_train, y_train)
    knnreg.fit(X_train, y_train)
    rafreg.fit(X_train, y_train)

    linreg_preds = linreg.predict(X_test)
    knnreg_preds = knnreg.predict(X_test)
    rafreg_preds = rafreg.predict(X_test)

    linreg_r2.append(metrics.r2_score(y_test, linreg_preds))
    linreg_mae.append(metrics.mean_absolute_error(y_test, linreg_preds))
    linreg_mse.append(metrics.mean_squared_error(y_test, linreg_preds))
    linreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪linreg_preds)),4))

    knnreg_r2.append(metrics.r2_score(y_test, knnreg_preds))
    knnreg_mae.append(metrics.mean_absolute_error(y_test, knnreg_preds))
    knnreg_mse.append(metrics.mean_squared_error(y_test, knnreg_preds))
    knnreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪knnreg_preds)),4))

```

```

    randfor_r2.append(metrics.r2_score(y_test, rafreg_preds))
    randfor_mae.append(metrics.mean_absolute_error(y_test, rafreg_preds))
    randfor_mse.append(metrics.mean_squared_error(y_test, rafreg_preds))
    randfor_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↳rafreg_preds)),4))

print("Linear Regression Average R^2: " + str(np.mean(linreg_r2)))
print("Linear Regression Average MAE: " + str(np.mean(linreg_mae)))
print("Linear Regression Average MSE: " + str(np.mean(linreg_mse)))
print("Linear Regression Average RMSE: " + str(np.mean(linreg_rmse)))

print("KNN Regression Average R^2: " + str(np.mean(knnreg_r2)))
print("KNN Regression Average MAE: " + str(np.mean(knnreg_mae)))
print("KNN Regression Average MSE: " + str(np.mean(knnreg_mse)))
print("KNN Regression Average RMSE: " + str(np.mean(knnreg_rmse)))

print("Random Forest Regression Average R^2: " + str(np.mean(randfor_r2)))
print("Random Forest Average MAE: " + str(np.mean(randfor_mae)))
print("Random Forest Average MSE: " + str(np.mean(randfor_mse)))
print("Random Forest Average RMSE: " + str(np.mean(randfor_rmse)))

```

```

Linear Regression Average R^2: 0.9561843546470922
Linear Regression Average MAE: 83589.4922107453
Linear Regression Average MSE: 123247494417.44443
Linear Regression Average RMSE: 295474.7414100001
KNN Regression Average R^2: 0.9429495156388995
KNN Regression Average MAE: 83783.73564100184
KNN Regression Average MSE: 140615550422.30243
KNN Regression Average RMSE: 335589.14285
Random Forest Regression Average R^2: 0.9542348536857073
Random Forest Average MAE: 79310.09288001957
Random Forest Average MSE: 125168736383.36545
Random Forest Average RMSE: 310490.96091

```

We find that applying PCA actually hurts the performance of our models so we do not employ it going forwards

```

[40]: # 11. Gridsearch to Optimize Parameters
#      Here we optimize the parameters on this dataset for the random forest
↳ensemble model we produced earlier
#      by default (which we employ above), the number of trees is 100, criterion
↳is "squared error", max_depth
#      is unlimited, bootstrap=True, and there are many others. The
↳hyperparameters were are going to experiment
#      with is the number of trees, the criterion, and max tree depth.
#

```

```

# Note: By default, GridSearchCV employs 5-fold cross validation

X_train, X_test, y_train, y_test = train_test_split(data_prepared, data_y,
    ↪train_size=0.85, random_state=122)

from sklearn.model_selection import GridSearchCV

param_grid = [
    {'n_estimators': [10, 50, 100, 200],
     'max_depth': [5, 10, 15]}
]

base_estimator = RandomForestRegressor()
sh = GridSearchCV(base_estimator, param_grid).fit(X_train, y_train)

sh.best_estimator_

```

[40]: RandomForestRegressor(max_depth=15, n_estimators=50)

[41]: *# implement hyperparameter tuned random forest*

```

optimal_randforest = RandomForestRegressor(max_depth=15, n_estimators=200)
optimal_randforest.fit(X_train, y_train)
opt_preds_train = optimal_randforest.predict(X_train)
opt_preds_test = optimal_randforest.predict(X_test)

print("Random Forest Metrics Following Hyperparameter Optimization")

print("\nTraining Set Metrics")
regression_results(y_train, opt_preds_train)
print("\nTesting Set Metrics")
regression_results(y_test, opt_preds_test)

```

Random Forest Metrics Following Hyperparameter Optimization

Training Set Metrics

```

explained_variance: 0.9913
r2: 0.9913
MAE: 46162.2107
MSE: 22770826332.4789
RMSE: 150900.0541

```

Testing Set Metrics

```

explained_variance: 0.98
r2: 0.98
MAE: 77355.783
MSE: 42009717731.5782

```

RMSE: 204962.7228

```
[42]: # 12. Experiment with Custom Model and Report Findings/Metrics

# Here we will try to train a neural network regressor to predict sales

from sklearn.neural_network import MLPRegressor

param_grid = [{
    'hidden_layer_sizes': [(10,10,10,10,10), (10,10,10,10,10,10,10),
                           (15,15,15,15,15), (15,15,15,15,15,15,15),
                           (20,20,20,20,20), (20,20,20,20,20,20,20)],
    'max_iter': [2000]
}]

base_estimator = MLPRegressor()
sh = GridSearchCV(base_estimator, param_grid).fit(X_train, y_train)

sh.best_estimator_
```

```
[42]: MLPRegressor(hidden_layer_sizes=(20, 20, 20, 20, 20), max_iter=2000)
```

```
[43]: nn = MLPRegressor(hidden_layer_sizes=(10,10,10,10,10), max_iter=500)
nn.fit(X_train, y_train)

nn_preds_train = nn.predict(X_train)
nn_preds_test = nn.predict(X_test)

print("Neural Network Regressor Metrics")

print("\nTraining Set Metrics")
regression_results(y_train, nn_preds_train)
print("\nTesting Set Metrics")
regression_results(y_test, nn_preds_test)
```

Neural Network Regressor Metrics Following PCA reduction to 12 components

Training Set Metrics

explained_variance: 0.9494

r2: 0.9494

MAE: 85570.7136

MSE: 132793678641.7084

RMSE: 364408.6698

Testing Set Metrics

explained_variance: 0.9825

r2: 0.9824

MAE: 83998.5607
MSE: 36921165101.2723
RMSE: 192148.8098

```
[45]: # Finalized Models and Performance Metrics Summary

from sklearn.metrics import mean_absolute_percentage_error

kfold = KFold(n_splits = 10, shuffle=True)
splits = kfold.split(data_prepared, data_y)

linreg_r2 = []
linreg_mae = []
linreg_mse = []
linreg_rmse = []
linreg_mape = []

knnreg_r2 = []
knnreg_mae = []
knnreg_mse = []
knnreg_rmse = []
knnreg_mape = []

randfor_r2 = []
randfor_mae = []
randfor_mse = []
randfor_rmse = []
randfor_mape = []

nnreg_r2 = []
nnreg_mae = []
nnreg_mse = []
nnreg_rmse = []
nnreg_mape = []

for train_indeces, test_indeces in splits:

    X_train = data_prepared[train_indeces]
    y_train = data_y.iloc[train_indeces]

    X_test = data_prepared[test_indeces, :]
    y_test = data_y.iloc[test_indeces]

    linreg = LinearRegression()
    knnreg = KNeighborsRegressor(n_neighbors = 3)
    rafreg = RandomForestRegressor(max_depth = 15, n_estimators = 50)
    nnreg = MLPRegressor(hidden_layer_sizes=(20, 20, 20, 20, 20), max_iter=2000)
```

```

linreg.fit(X_train, y_train)
knnreg.fit(X_train, y_train)
rafreg.fit(X_train, y_train)
nnreg.fit(X_train, y_train)

linreg_preds = linreg.predict(X_test)
knnreg_preds = knnreg.predict(X_test)
rafreg_preds = rafreg.predict(X_test)
nnreg_preds = nnreg.predict(X_test)

linreg_r2.append(metrics.r2_score(y_test, linreg_preds))
linreg_mae.append(metrics.mean_absolute_error(y_test, linreg_preds))
linreg_mse.append(metrics.mean_squared_error(y_test, linreg_preds))
linreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪linreg_preds)),4))
linreg_mape.append(mean_absolute_percentage_error(y_test, linreg_preds))

knnreg_r2.append(metrics.r2_score(y_test, knnreg_preds))
knnreg_mae.append(metrics.mean_absolute_error(y_test, knnreg_preds))
knnreg_mse.append(metrics.mean_squared_error(y_test, knnreg_preds))
knnreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪knnreg_preds)),4))
knnreg_mape.append(mean_absolute_percentage_error(y_test, knnreg_preds))

randfor_r2.append(metrics.r2_score(y_test, rafreg_preds))
randfor_mae.append(metrics.mean_absolute_error(y_test, rafreg_preds))
randfor_mse.append(metrics.mean_squared_error(y_test, rafreg_preds))
randfor_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪rafreg_preds)),4))
randfor_mape.append(mean_absolute_percentage_error(y_test, rafreg_preds))

nnreg_r2.append(metrics.r2_score(y_test, nnreg_preds))
nnreg_mae.append(metrics.mean_absolute_error(y_test, nnreg_preds))
nnreg_mse.append(metrics.mean_squared_error(y_test, nnreg_preds))
nnreg_rmse.append(round(np.sqrt(metrics.mean_squared_error(y_test,
↪nnreg_preds)),4))
nnreg_mape.append(mean_absolute_percentage_error(y_test, nnreg_preds))

print("Linear Regression")
print("R^2: " + str(np.mean(linreg_r2)))
print("MAE: " + str(np.mean(linreg_mae)))
print("MSE: " + str(np.mean(linreg_mse)))
print("RMSE: " + str(np.mean(linreg_rmse)))
print("MAPE: " + str(np.mean(linreg_mape)))

```

```

print("\nKNN Regression")
print("R^2: " + str(np.mean(knnreg_r2)))
print("MAE: " + str(np.mean(knnreg_mae)))
print("MSE: " + str(np.mean(knnreg_mse)))
print("RMSE: " + str(np.mean(knnreg_rmse)))
print("MAPE: " + str(np.mean(knnreg_mape)))

print("\nRandom Forest Regression")
print("R^2: " + str(np.mean(randfor_r2)))
print("MAE: " + str(np.mean(randfor_mae)))
print("MSE: " + str(np.mean(randfor_mse)))
print("RMSE: " + str(np.mean(randfor_rmse)))
print("MAPE: " + str(np.mean(randfor_mape)))

print("\nNeural Network Regression")
print("R^2: " + str(np.mean(nnreg_r2)))
print("MAE: " + str(np.mean(nnreg_mae)))
print("MSE: " + str(np.mean(nnreg_mse)))
print("RMSE: " + str(np.mean(nnreg_rmse)))
print("MAPE: " + str(np.mean(nnreg_mape)))

```

Linear Regression

R²: 0.9506262229145209
MAE: 83673.11524785537
MSE: 123681722060.72571
RMSE: 294337.1583
MAPE: 4.891075910234255e+16

KNN Regression

R²: 0.9493492344416433
MAE: 80769.3875125602
MSE: 126561893173.7651
RMSE: 312018.23373000004
MAPE: 8148317432747818.0

Random Forest Regression

R²: 0.93763077495652
MAE: 78643.25576040288
MSE: 136881793505.19601
RMSE: 321320.1806
MAPE: 2.4269400126400844e+16

Neural Network Regression

R²: 0.9521027275545263
MAE: 83010.9176332785
MSE: 120426132575.03174
RMSE: 289853.4439
MAPE: 5.3730673640729416e+16

[]: